

# Computer Vision

DATA.ML.300, 5 study credits

Esa Rahtu  
Unit of Computing Sciences, Tampere University

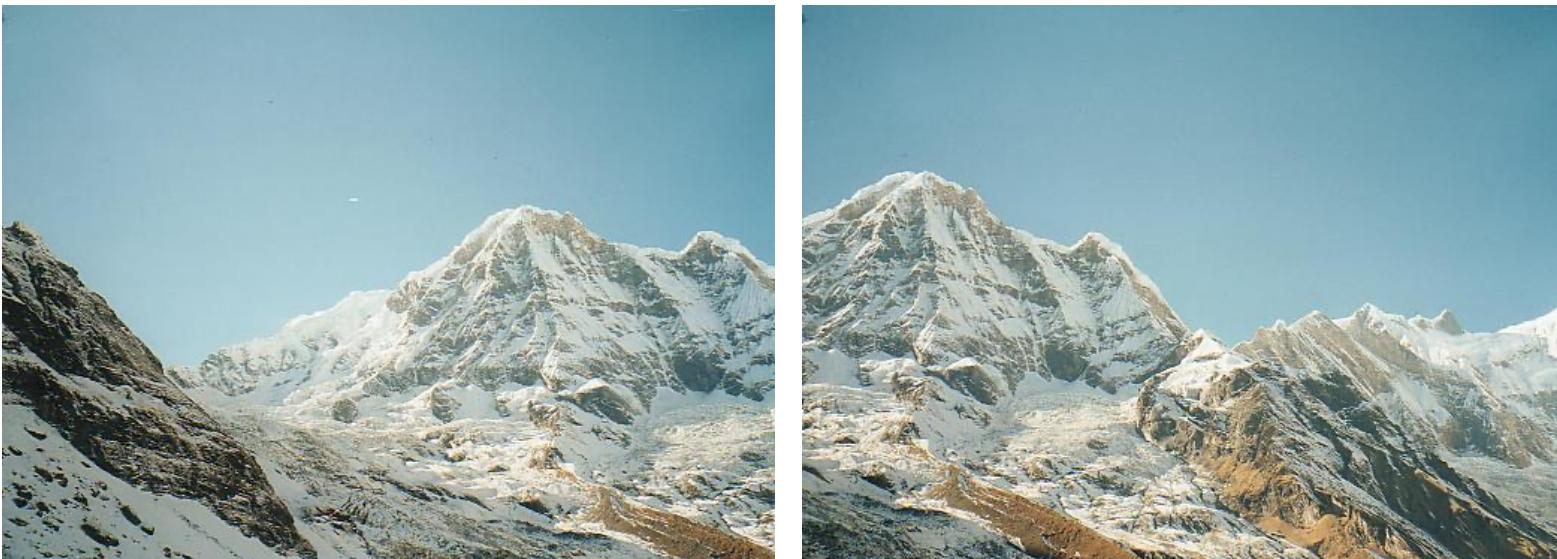
# Local image features

- Low-level features (corners, edges, texture patches) are needed for many tasks:
  - Image matching and registration
  - Structure-from-motion and image-based 3D modeling
  - Image segmentation
  - Object recognition
  - Image retrieval
- Relevant reading:
  - Chapter 4 of Szeliski's book
  - David Lowe's article (2004) <http://www.cs.ubc.ca/~lowe/keypoints/>

Acknowledgement: many slides from Svetlana Lazebnik, Steve Seitz, David Lowe, Kristen Grauman, and others (detailed credits on individual slides)

# Why to extract local features?

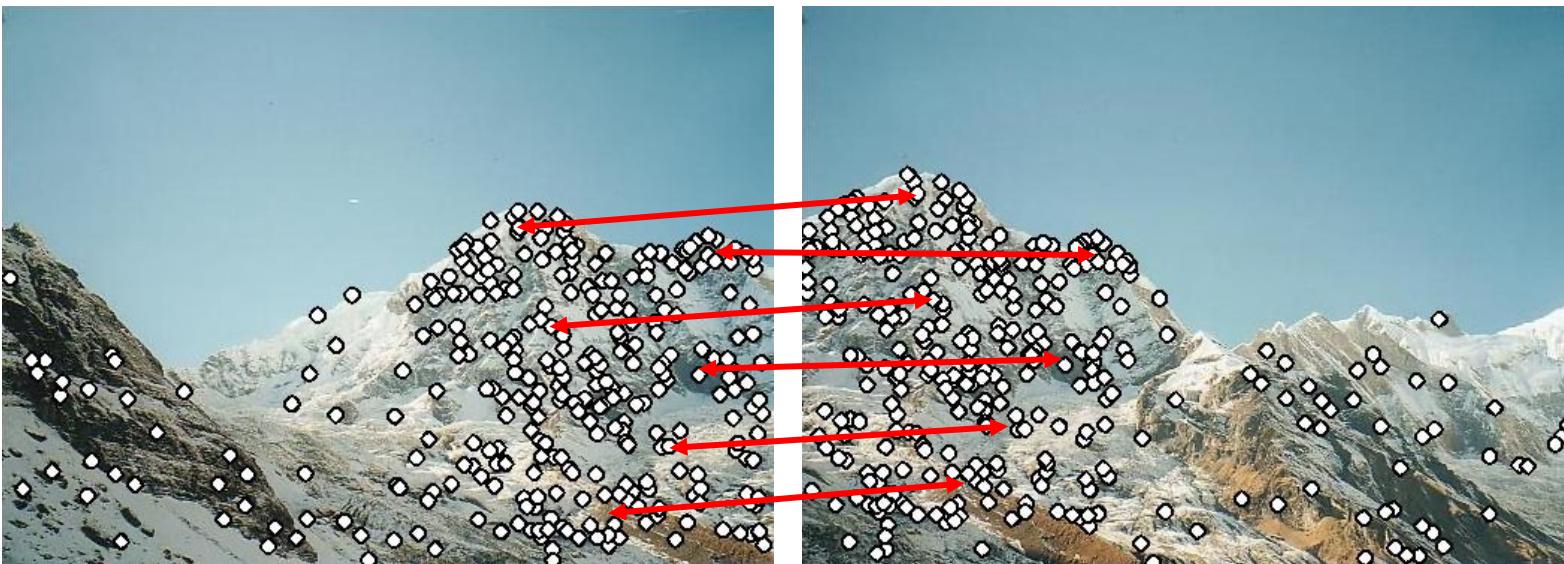
- Motivation: panorama stitching
  - We have two images – how do we combine them?



Source: S. Lazebnik

# Why to extract local features?

- Motivation: panorama stitching
  - We have two images – how do we combine them?



Step 1: extract local features (e.g. keypoints)

Step 2: find correspondences (e.g. match keypoint features)

# Why to extract local features?

- Motivation: panorama stitching
  - We have two images – how do we combine them?

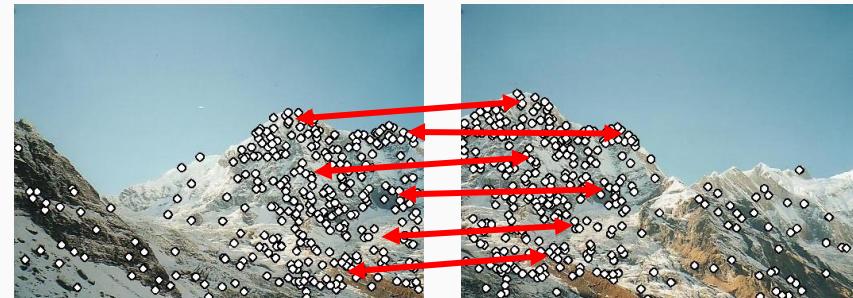
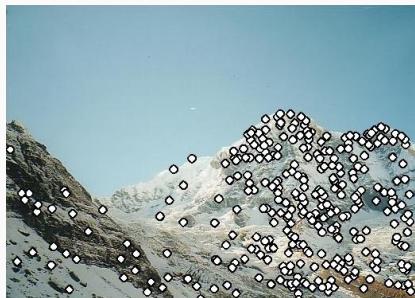
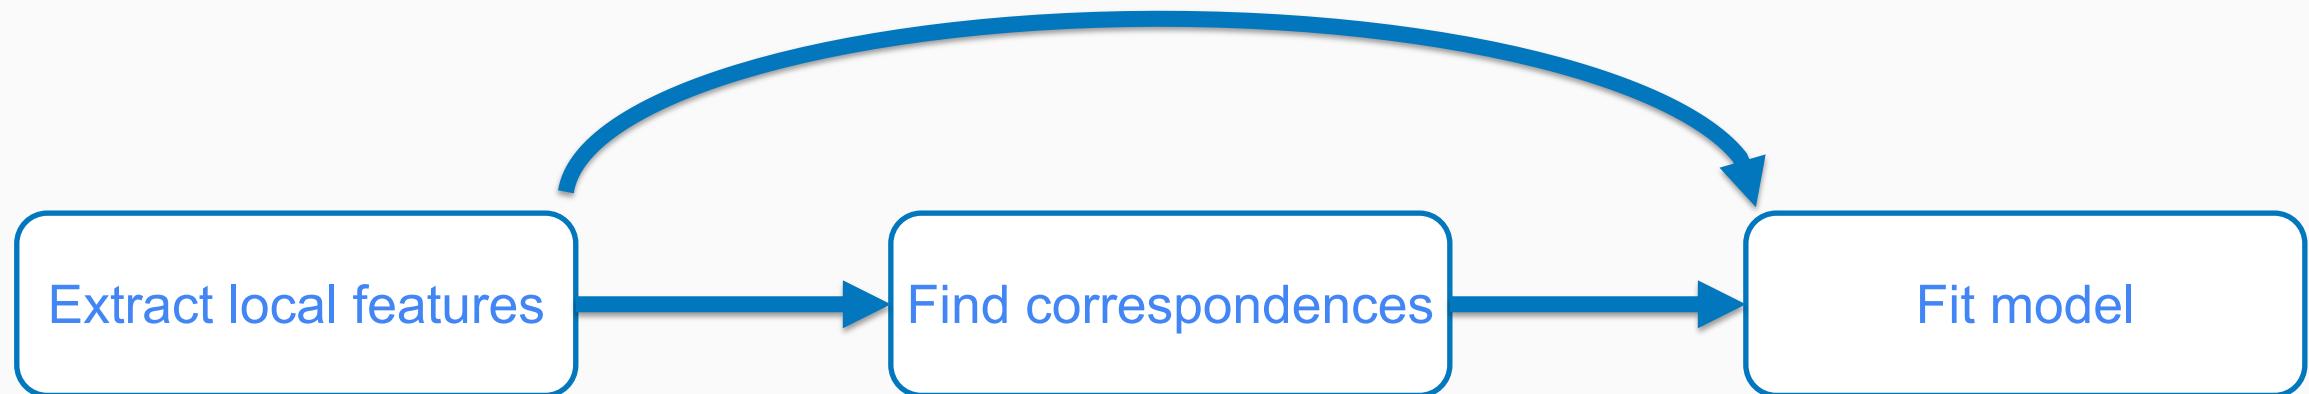


Step 1: extract local features (e.g. keypoints)

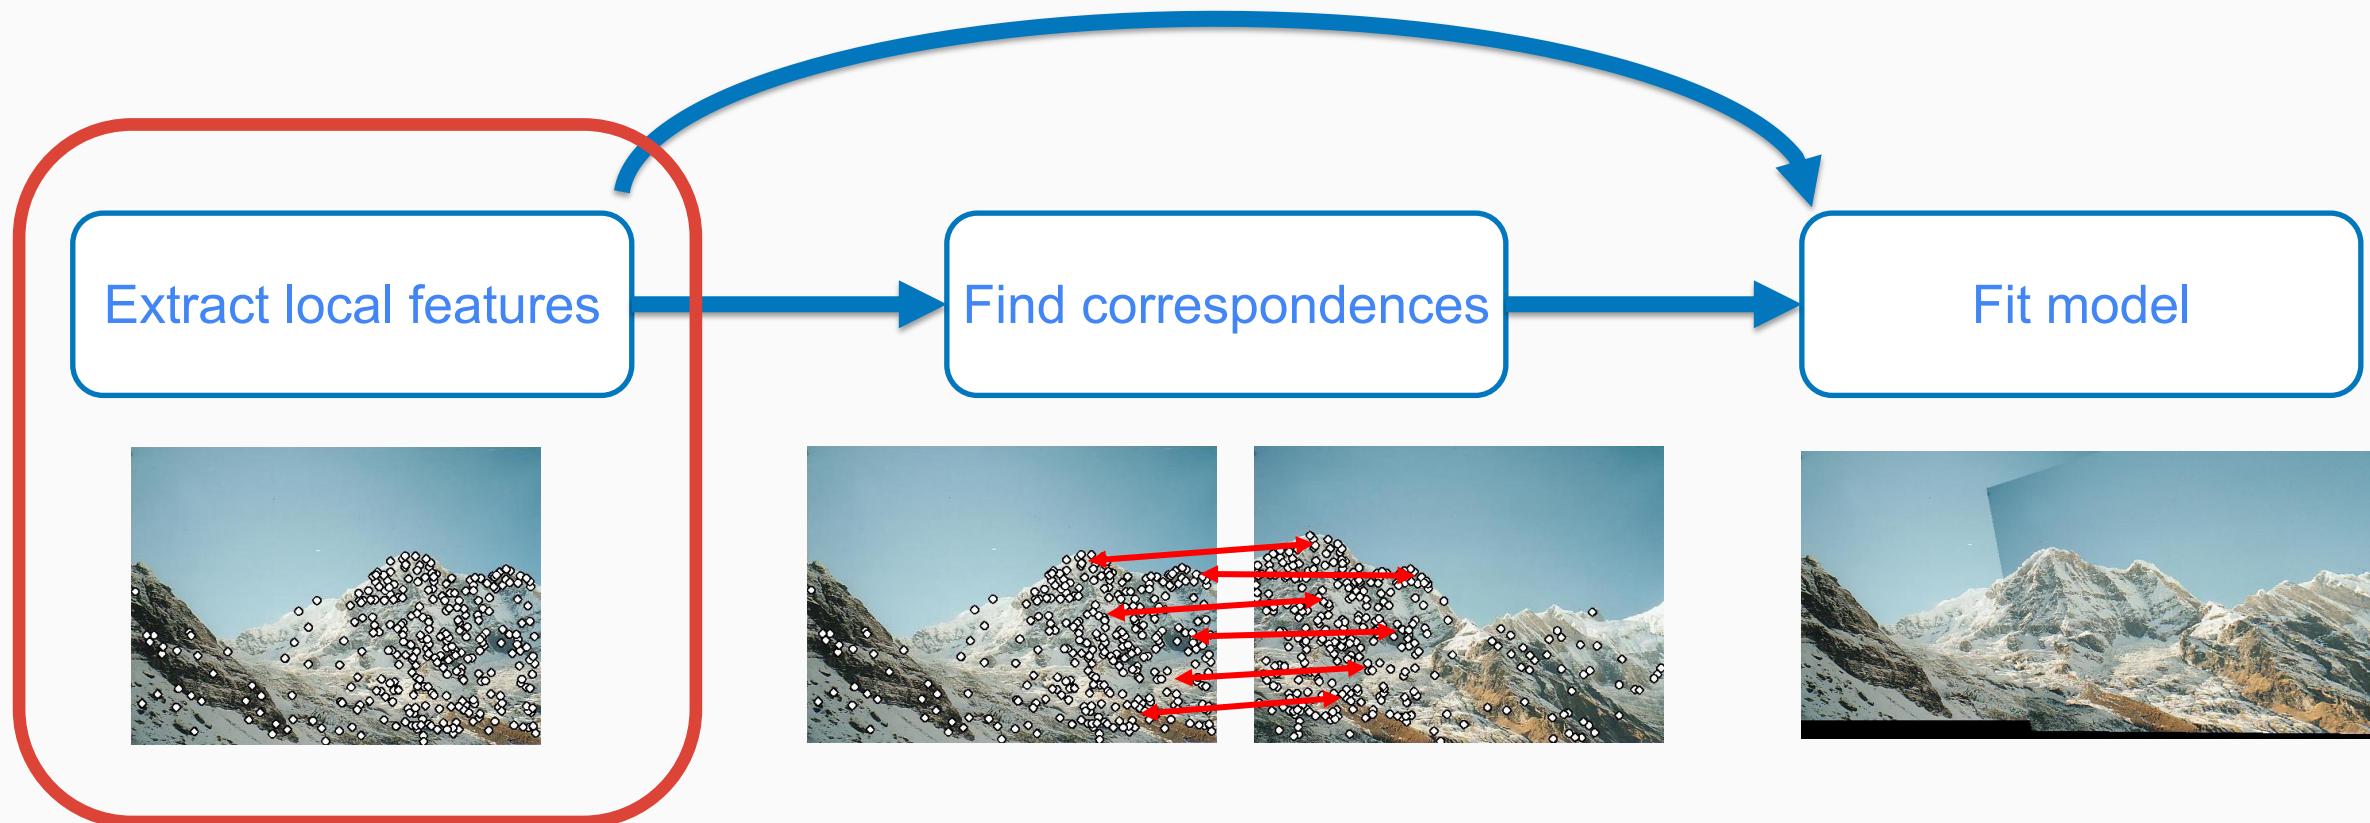
Step 2: find correspondences (e.g. match keypoint features)

Step 3: align images (e.g. model fitting)

# Usual steps in local feature extraction and matching



# Usual steps in local feature extraction and matching



# Edges



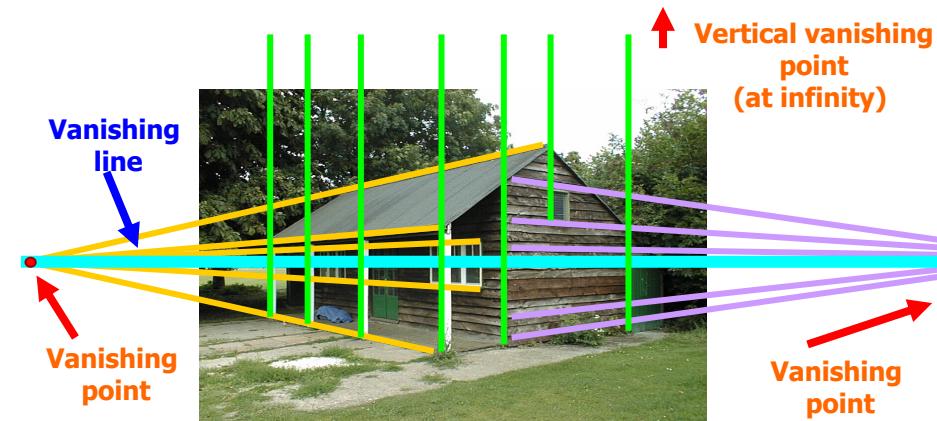
# Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image
  - Intuitively, most semantic and shape information from the image can be encoded in the edges
  - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)



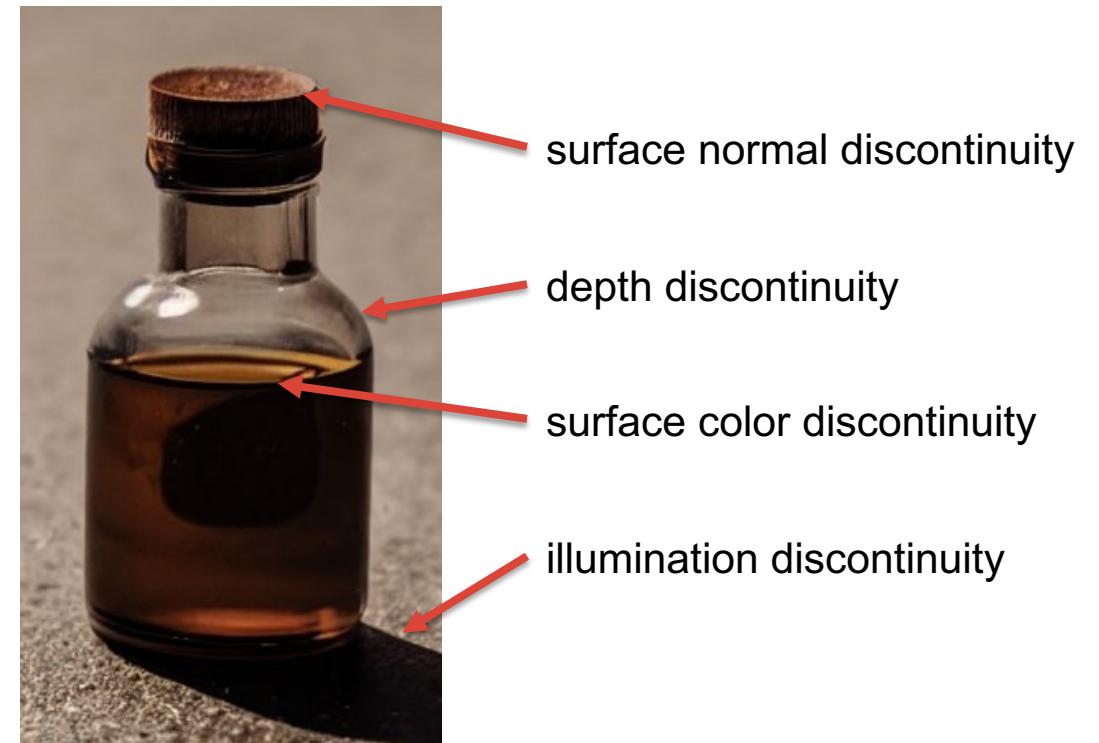
# Why do we care about edges?

- Extract information, recognize objects
- Recover geometry and viewpoint



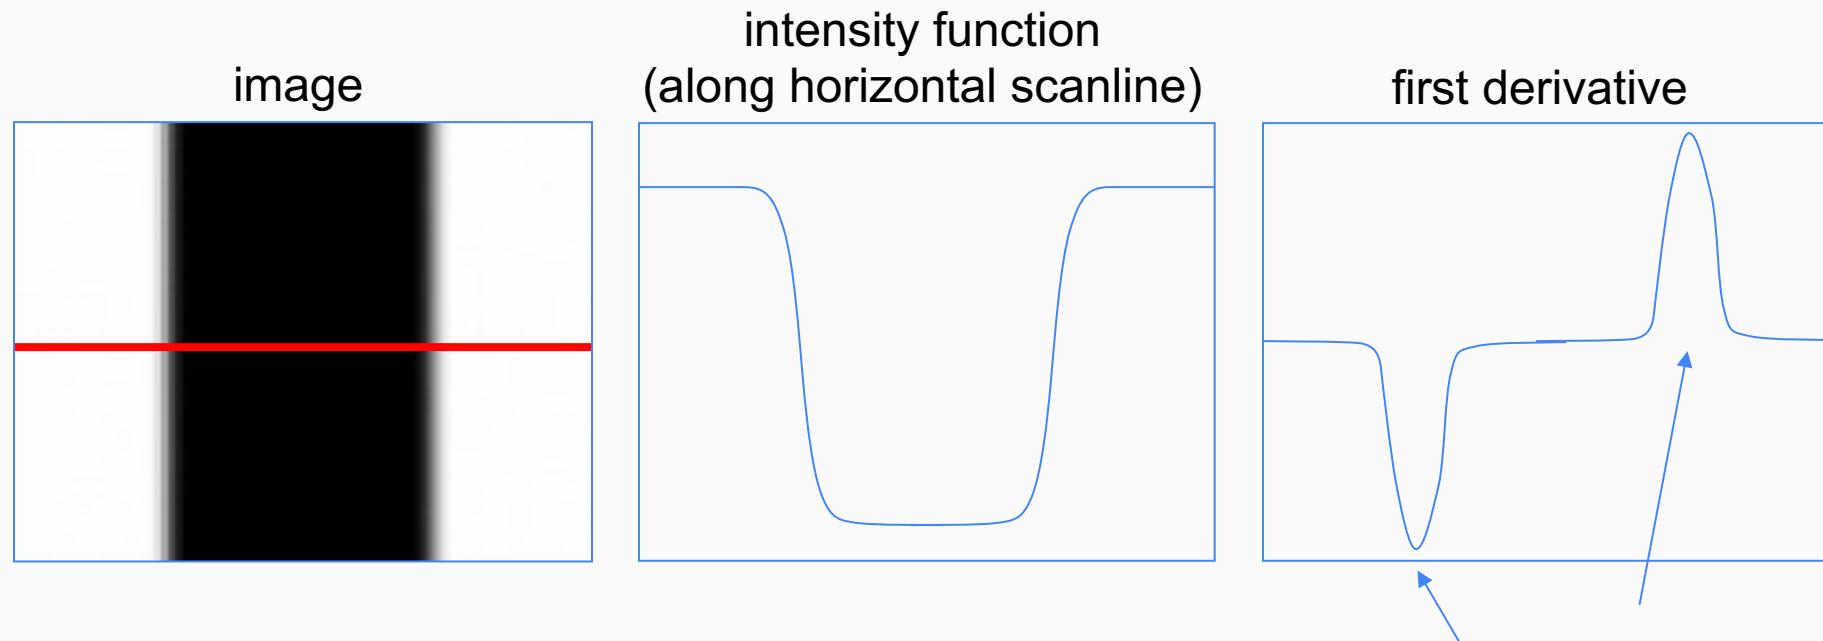
# Origin of edges

- Edges are caused by a variety of factors:



# Edge detection

- An edge is a place of rapid change in the image intensity function



edges correspond to  
extrema of derivative

# Reminder: Derivatives with convolution

- For 2D function  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

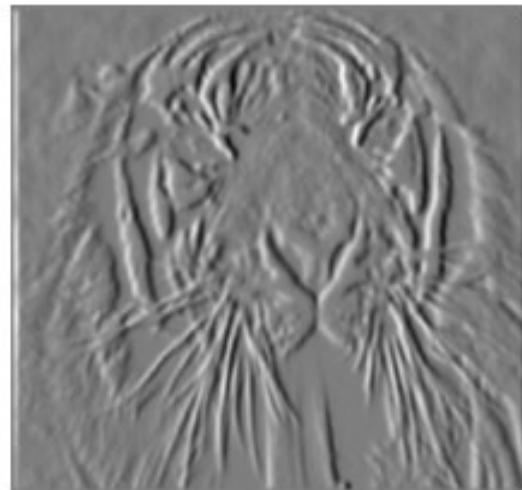
$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement the above as convolution, what would be the associated filter?

# Reminder: Partial derivatives of an image

$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

1
-1

## Reminder: Finite difference filters

- Other approximations of derivative filters exist:

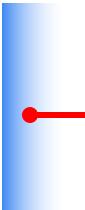
**Prewitt:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

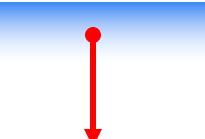
**Sobel:**  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

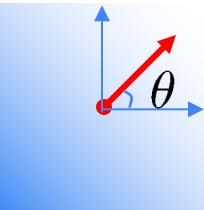
**Roberts:**  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

# Image gradient

The gradient of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid increase in intensity

- How does this direction relate to the direction of the edge?

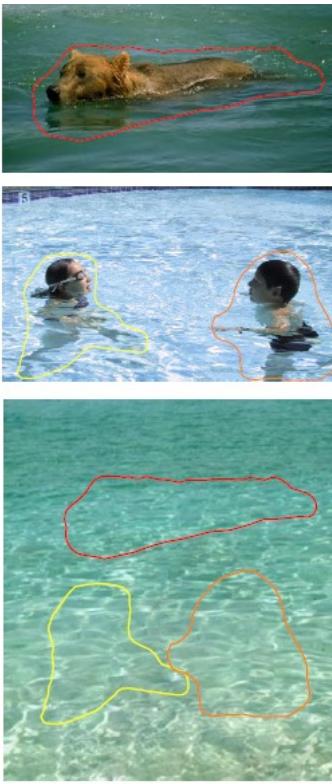
The gradient direction is given by  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

The edge strength is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Application: Gradient-domain image editing

- Goal: solve for pixel values in the target region to match gradients of the source region while keeping background pixels the same



sources/destinations



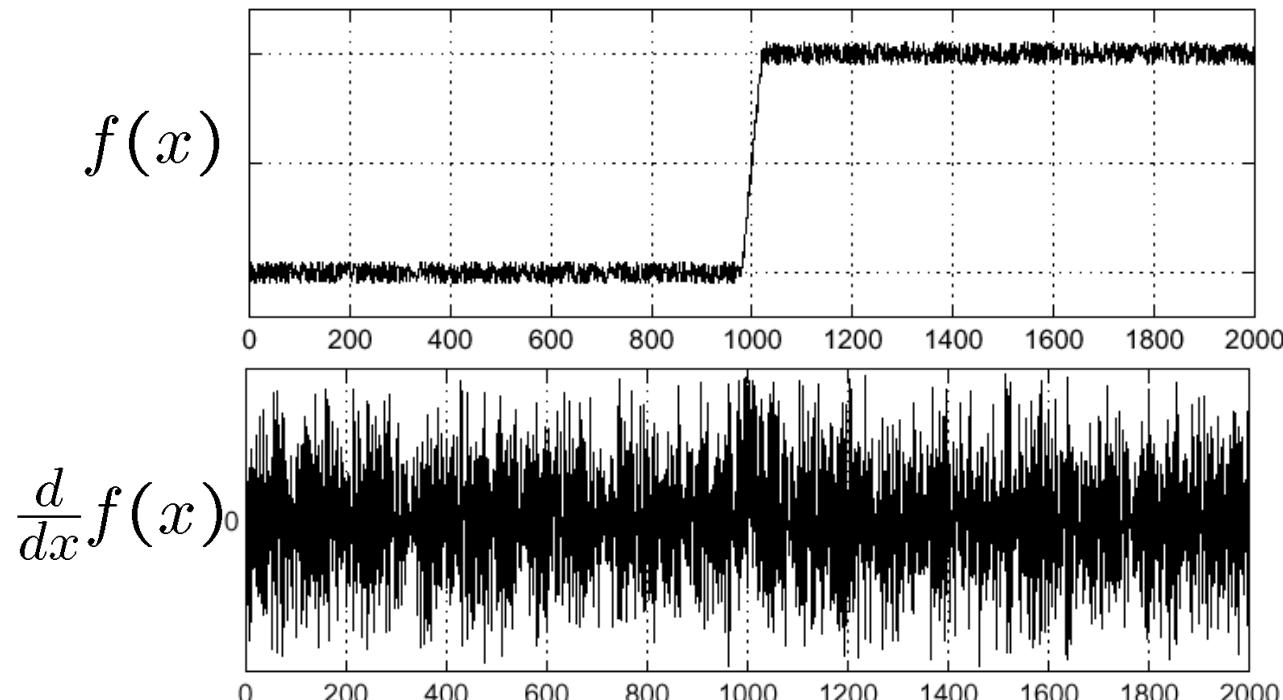
cloning



seamless cloning

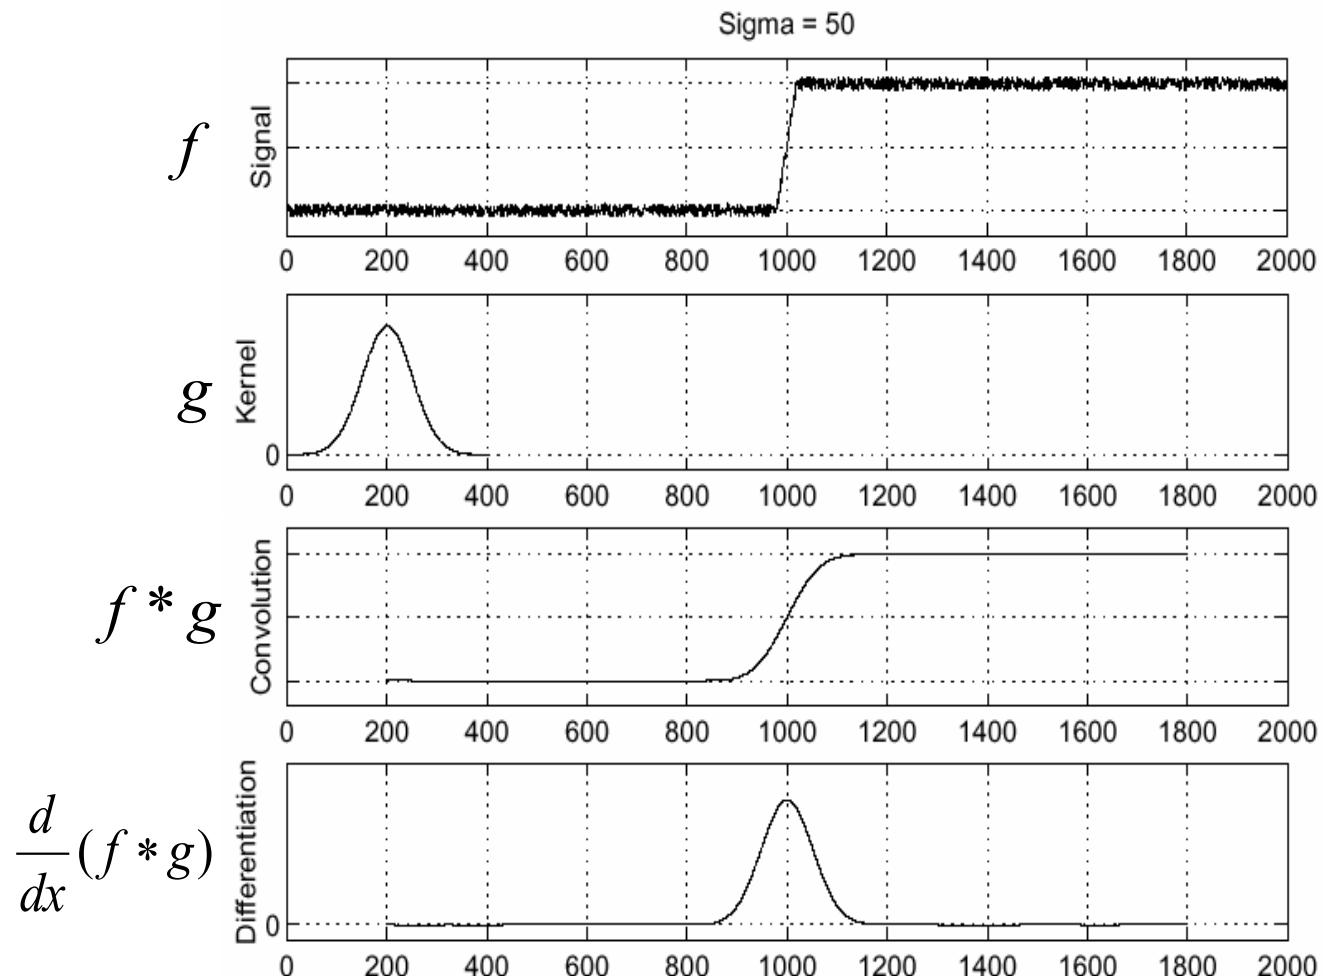
# Effects of noise

- Consider a single row or column of the image



Where is the edge?

Solution: smooth first



To find edges, look for peaks in

$$\frac{d}{dx}(f * g)$$

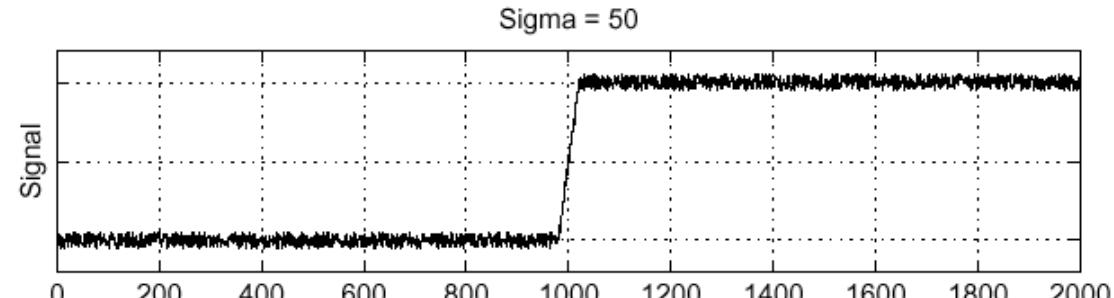
Source: S. Seitz

# Derivative theorem of convolution

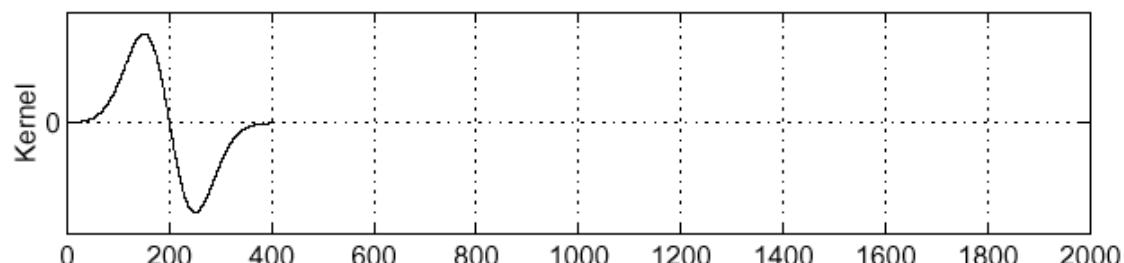
- Differentiation is convolution, and convolution is associative:
- This saves us one operation:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

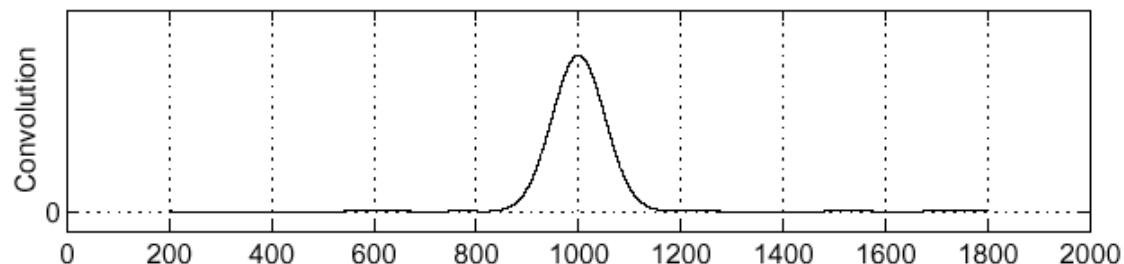
$f$



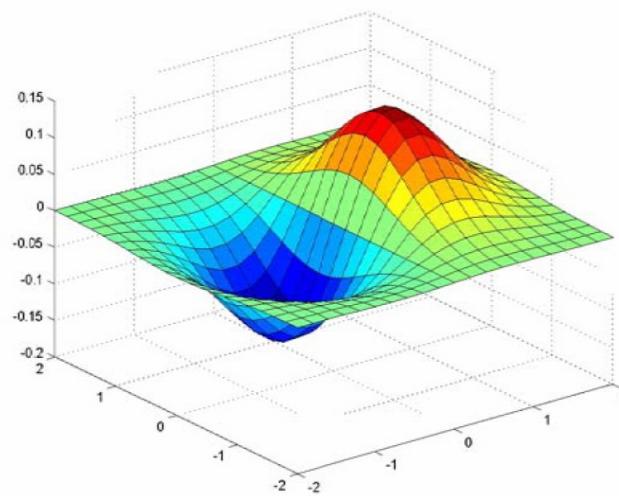
$\frac{d}{dx} g$



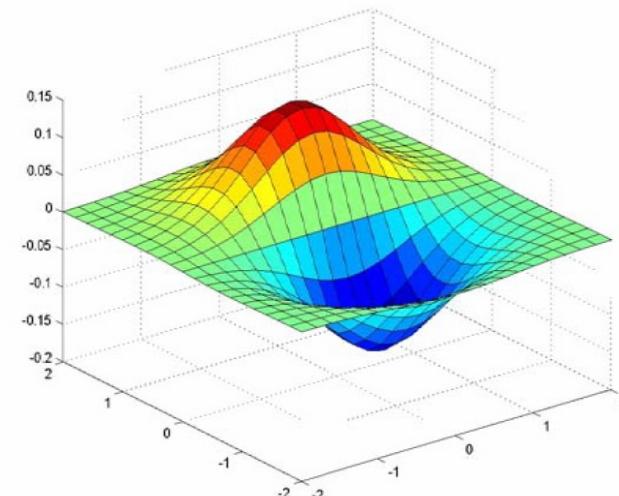
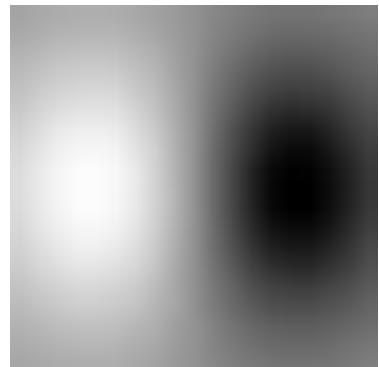
$f * \frac{d}{dx} g$



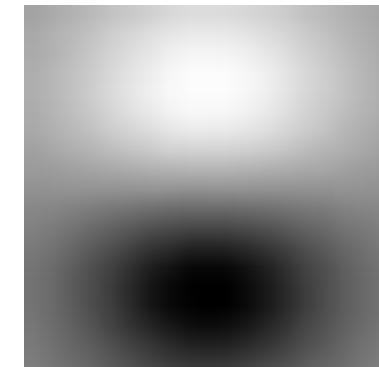
# Derivative of Gaussian filters



x-direction

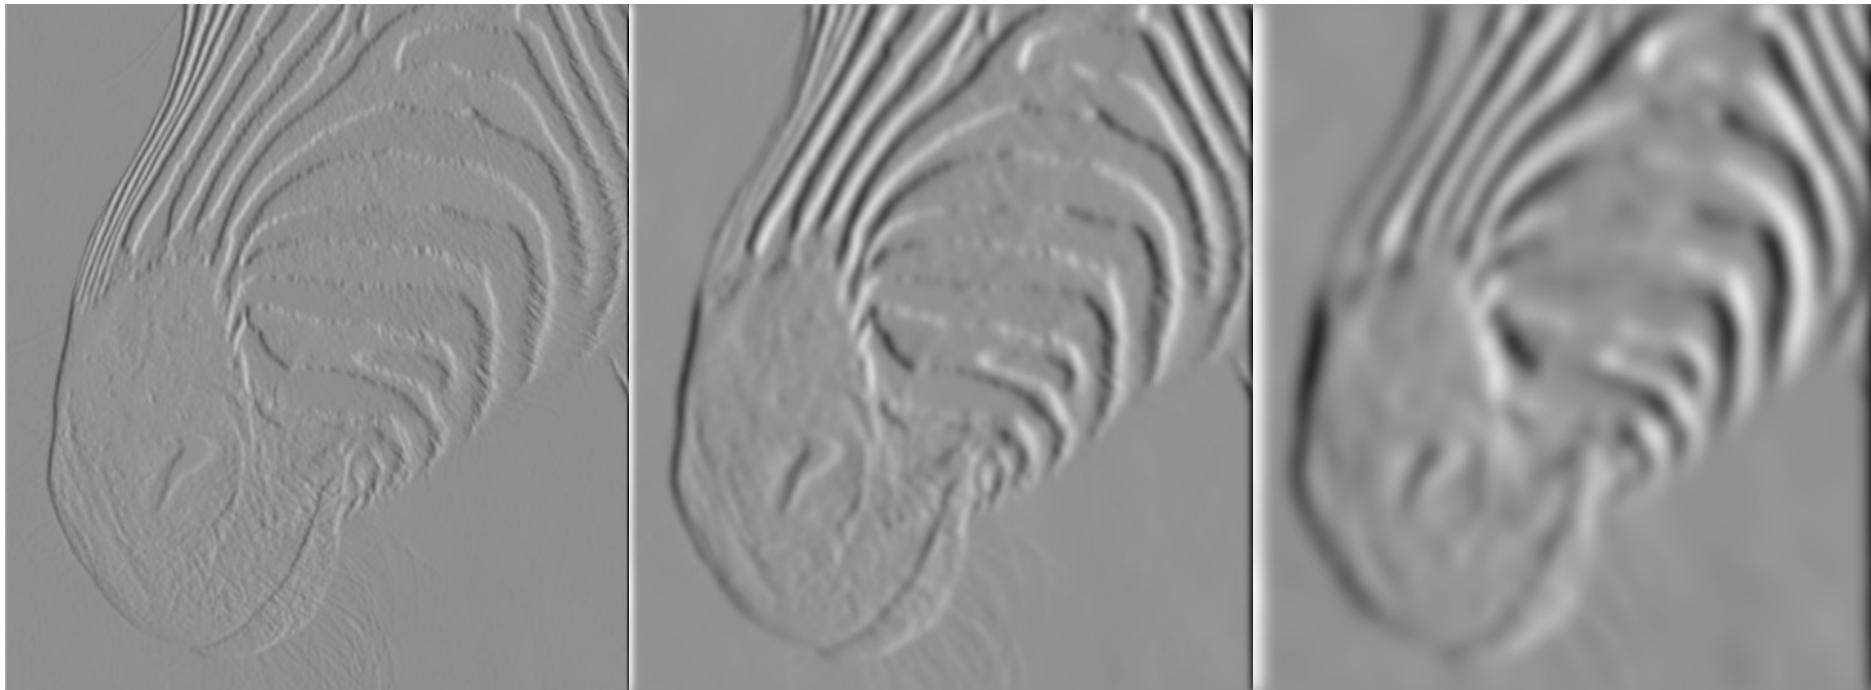


y-direction



# Scale of Gaussian derivative filter

- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”



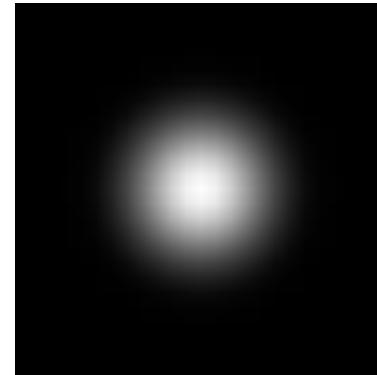
1 pixel

3 pixels

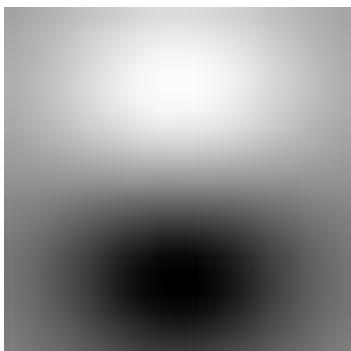
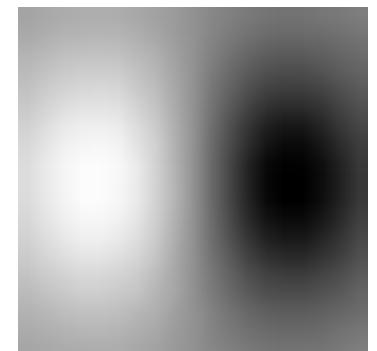
7 pixels

# Review: Smoothing vs. derivative filters

- Smoothing filters
  - Gaussian: remove “high-frequency” components; “low-pass” filter
  - Can the values of a smoothing filter be negative?
  - What should the values sum to?  
**One:** constant regions are not affected by the filter



- Derivative filters
  - Derivatives of Gaussian
  - Can the values of a derivative filter be negative?
  - What should the values sum to?  
**Zero:** no response in constant regions



# Building an edge detector



Original image



Final output

# Building an edge detector



Norm of the gradient

Source: S. Lazebnik

# Building an edge detector

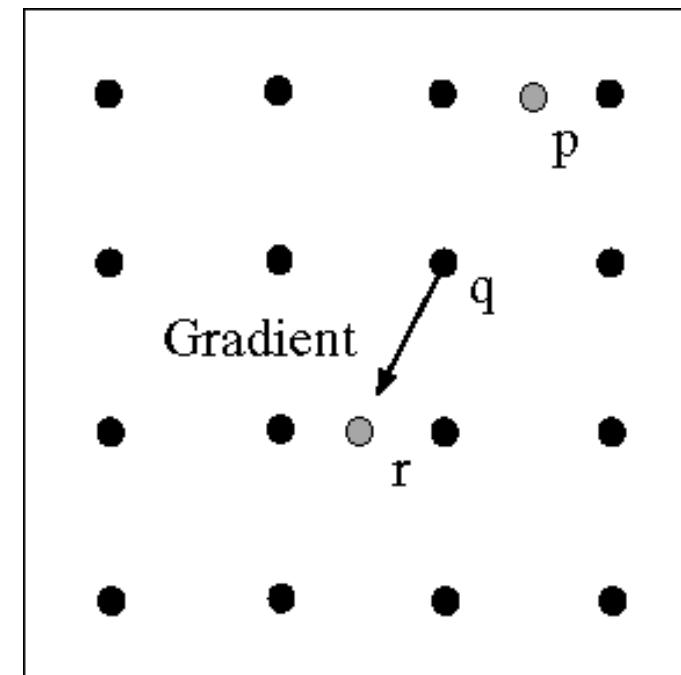
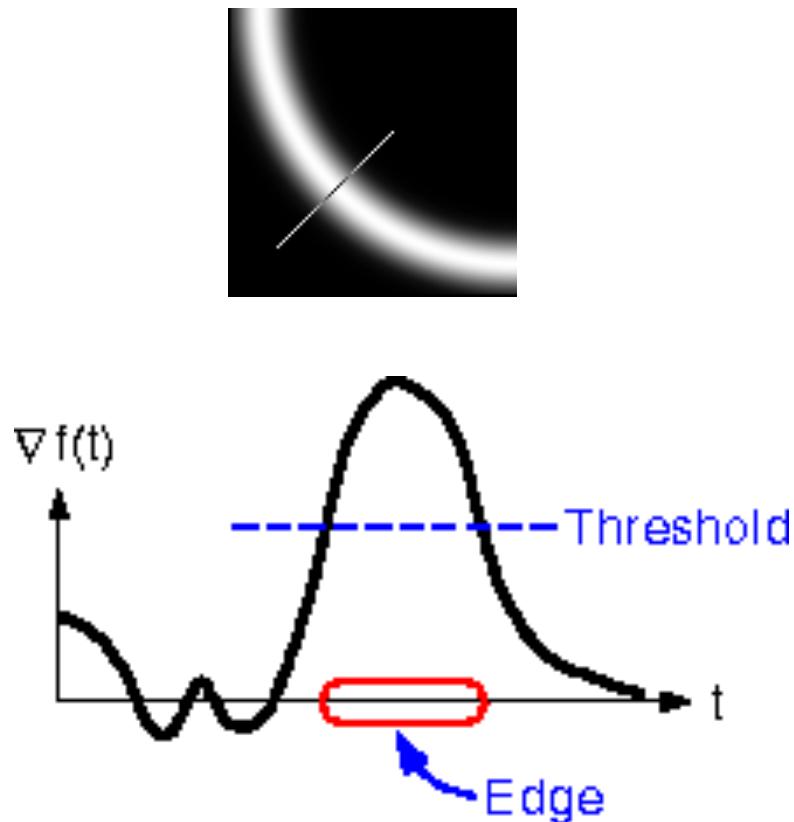


Thresholded norm of the gradient

How to turn  
these thick  
regions of  
the  
gradient  
into  
curves?

# Non-maximum suppression

- Check if pixel is local maximum along gradient direction, select single max across width of the edge
  - requires checking interpolated pixels p and r



# Non-maximum suppression



Another  
problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding

# Hysteresis thresholding

- Use a high threshold to start edge curves, and a low threshold to continue them.



# Hysteresis thresholding



original image



high threshold  
(strong edges)



low threshold  
(weak edges)



hysteresis threshold

Source: L. Fei-Fei

## Recap: Canny edge detector

1. Compute x and y gradient images
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin wide “ridges” down to single pixel width
4. Linking and thresholding (hysteresis):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them

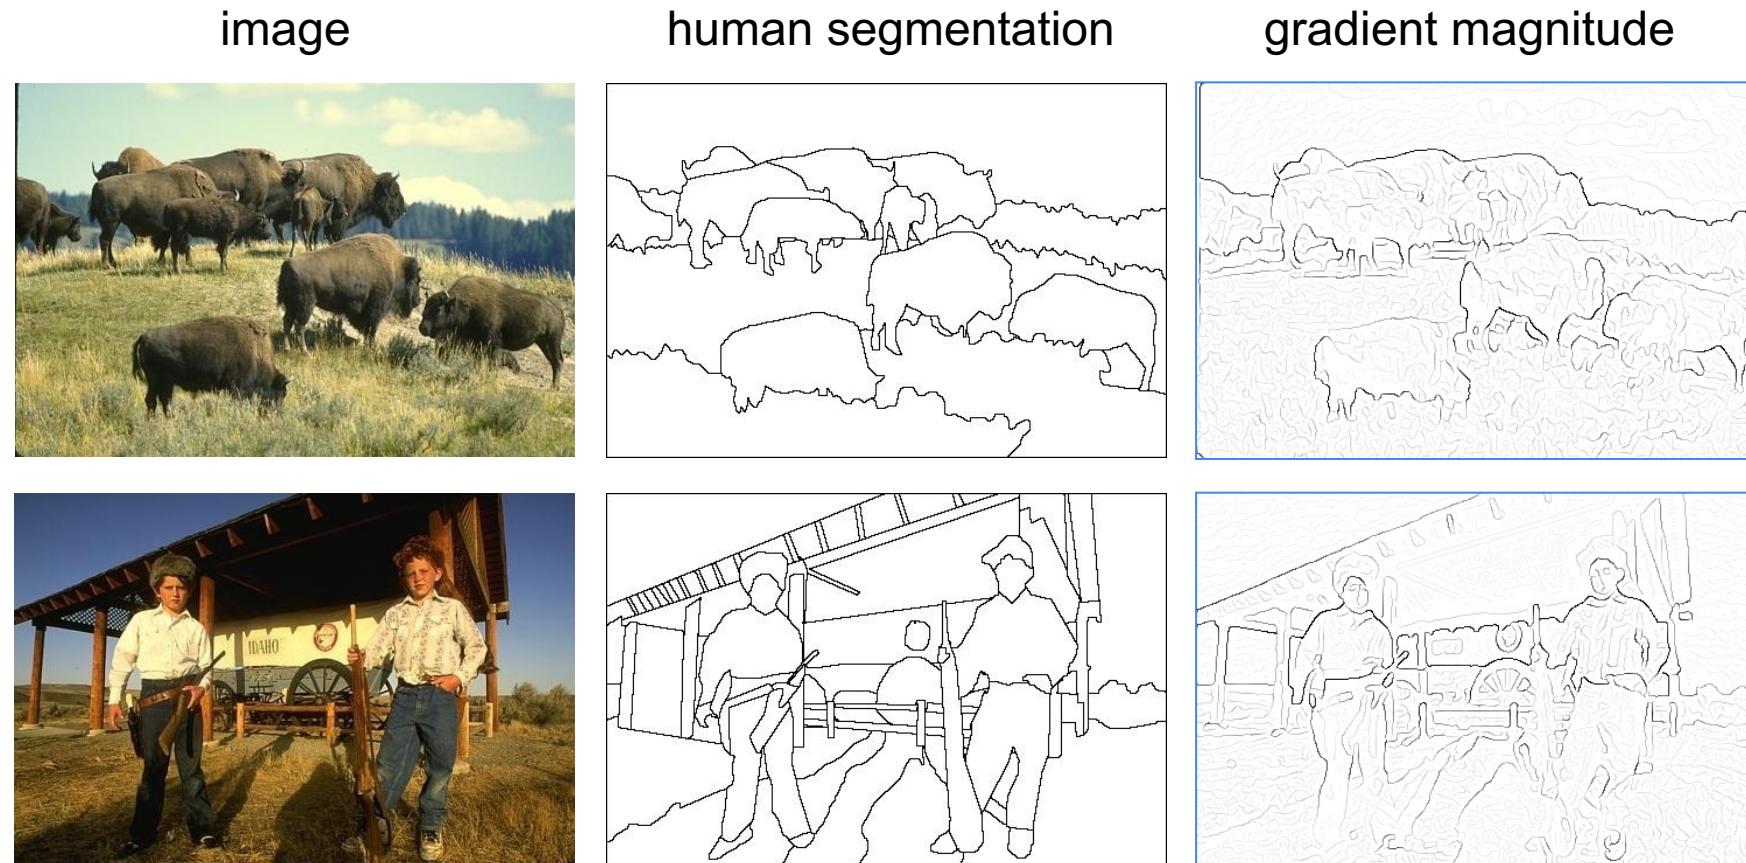
MATLAB: **edge(image, 'canny');**

J. Canny, [\*\*\*A Computational Approach To Edge Detection\*\*\*](#), IEEE  
Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

# Image gradients vs. meaningful contours

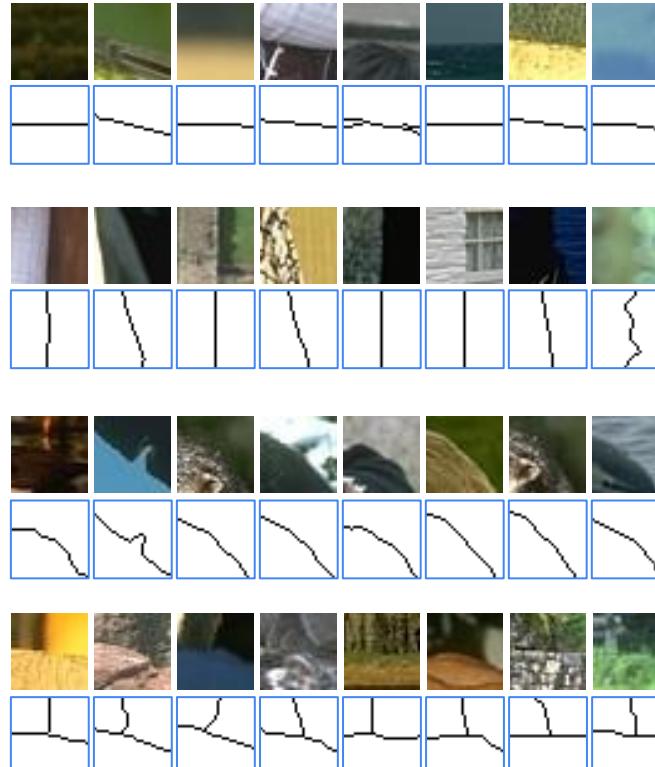
- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>



# Data-driven edge detection

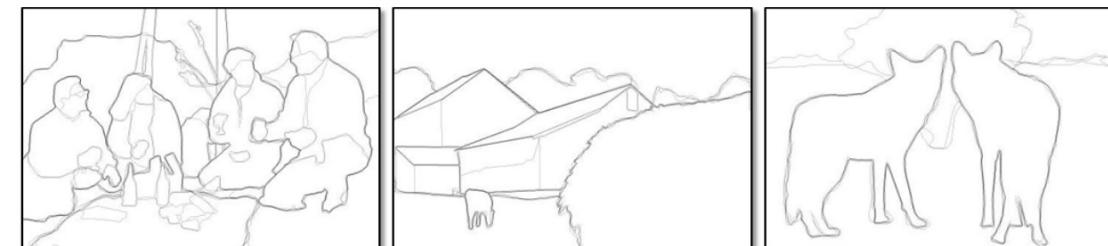
Training data



Input images



Ground truth



Output



# Keypoints



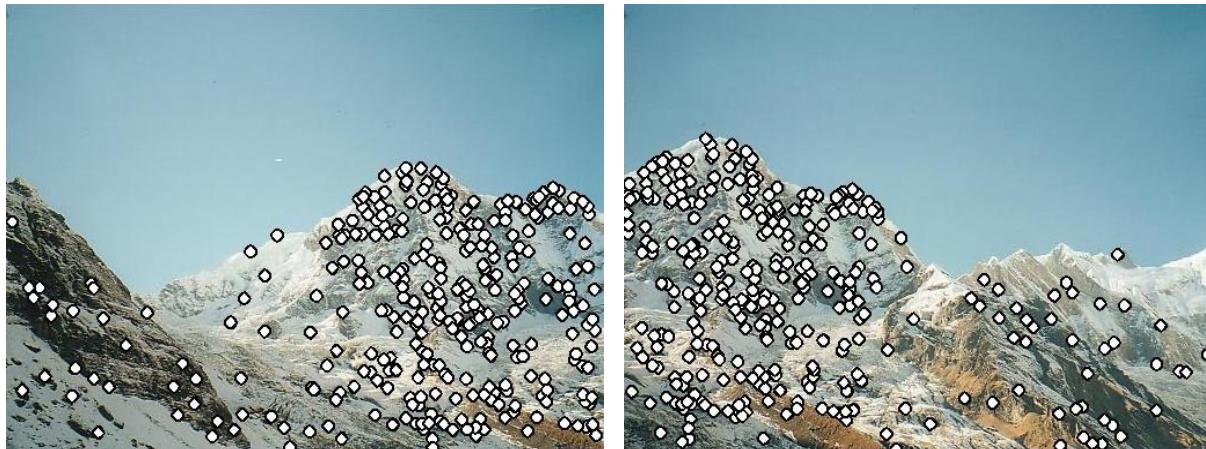
# Keypoint extraction: Corners



Source: S. Lazebnik

# Reminder: Panorama stitching

- Motivation: panorama stitching
  - We have two images – how do we combine them?



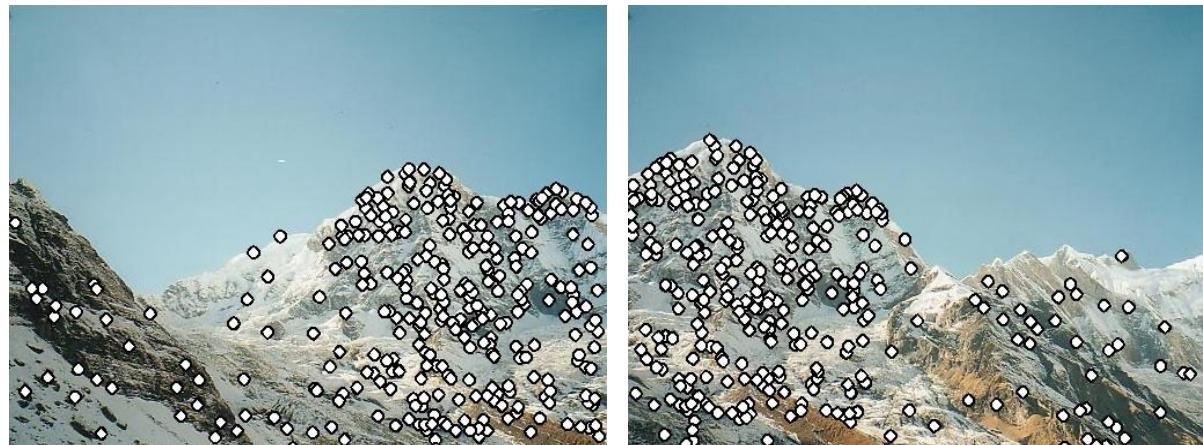
Step 1: extract keypoints

Step 2: match keypoint features

Step 3: align images

# Characteristics of good keypoints

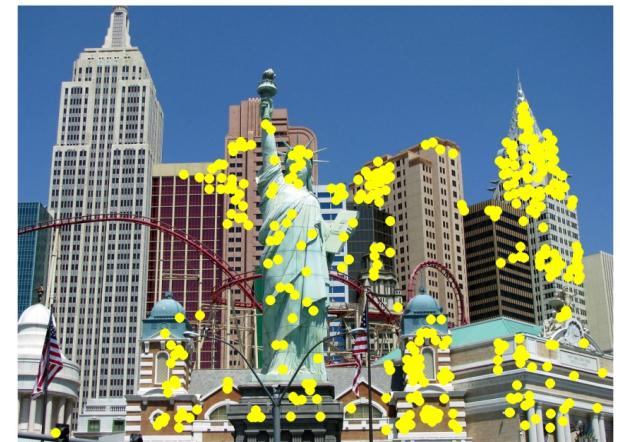
- Repeatability
  - The same keypoint can be found in several images despite geometric and photometric transformations
- Saliency
  - Each keypoint is distinctive
- Compactness and efficiency
  - Many fewer keypoints than image pixels
- Locality
  - A keypoint occupies a relatively small area of the image; robust to clutter and occlusion



Source: S. Lazebnik

# Applications

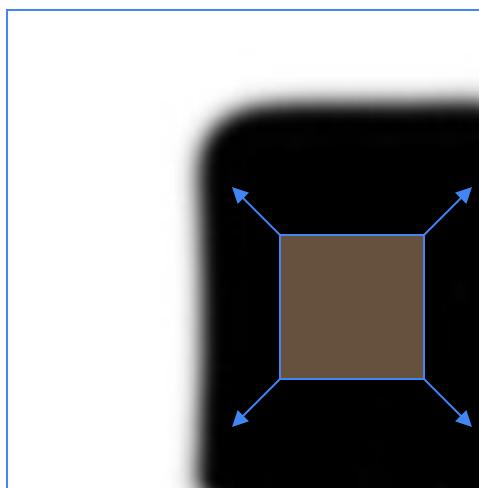
- Keypoints are used for:
  - Image alignment
  - 3D reconstruction
  - Motion tracking
  - Robot navigation
  - Indexing and database retrieval
  - Object recognition



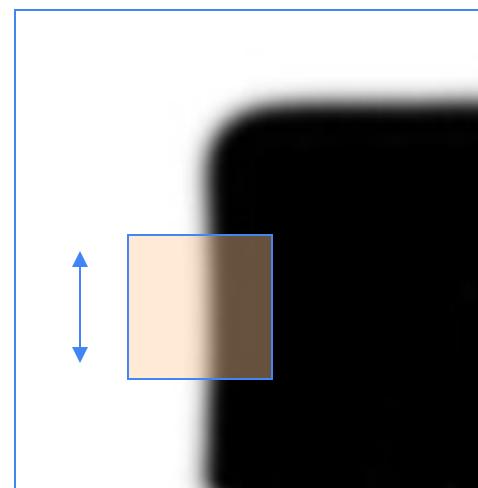
Source: S. Lazebnik

# Corner Detection: Basic Idea

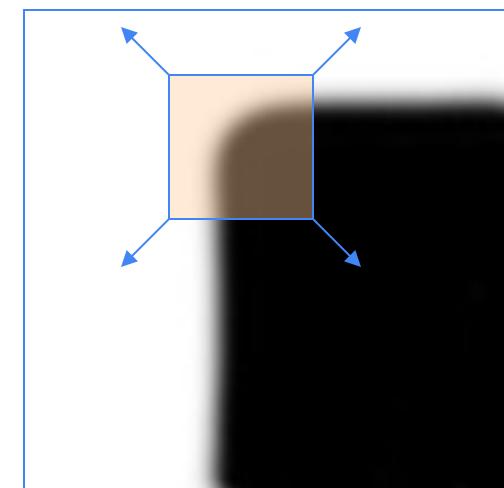
- We should easily recognize the point by looking through a small window
- Shifting a window in any direction should give a large change in intensity



**“flat” region:**  
no change in  
all directions



**“edge”:**  
no change  
along the edge  
direction

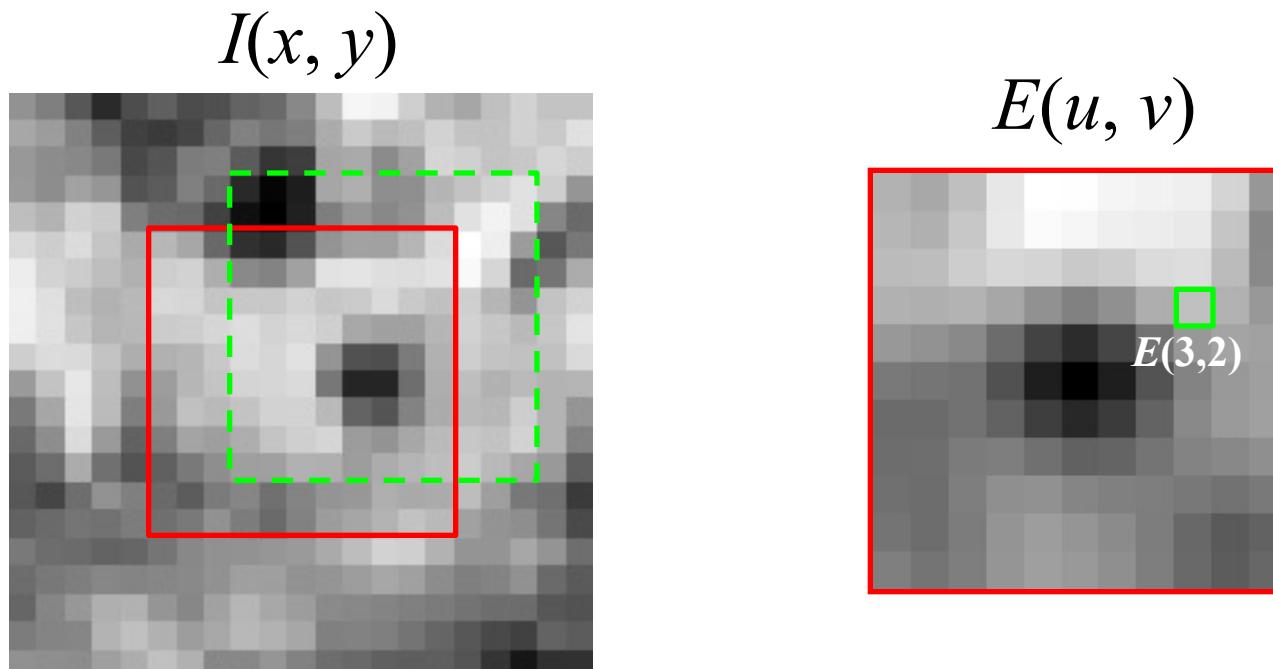


**“corner”:**  
significant  
change in all  
directions

# Corner Detection: Mathematics

Change in appearance of window W for the shift  $[u,v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

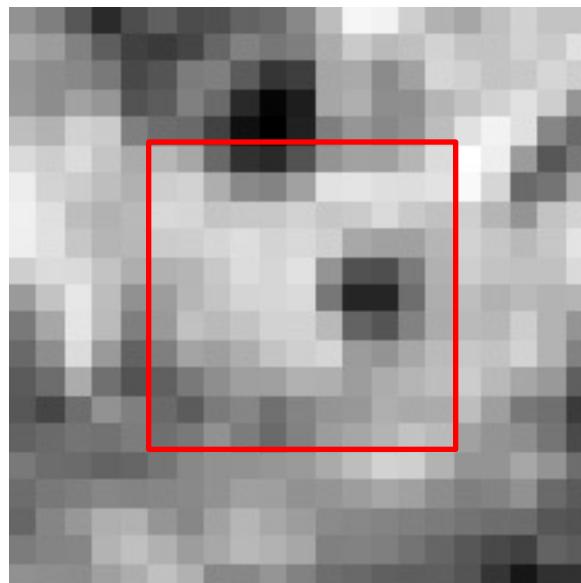


# Corner Detection: Mathematics

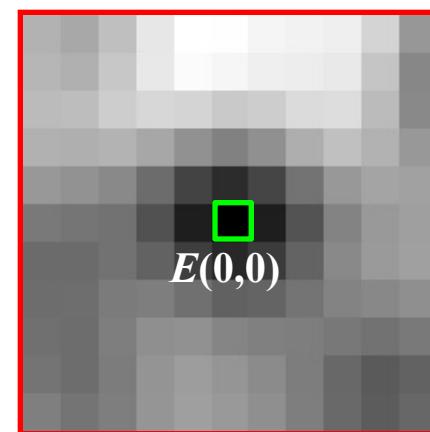
Change in appearance of window W for the shift  $[u,v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



$E(u, v)$



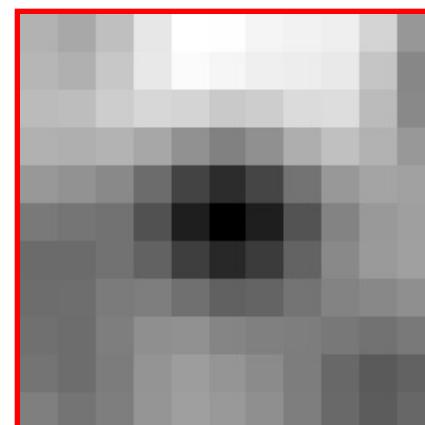
# Corner Detection: Mathematics

Change in appearance of window W for the shift  $[u,v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$



# Corner Detection: Mathematics

- First-order Taylor approximation for small motions  $[u, v]$ :

$$I(x+u, y+v) \approx I(x, y) + I_x u + I_y v$$

- Let's plug this into  $E(u, v)$ :

$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x+u, y+v) - I(x, y)]^2 \\ &\approx \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 \\ &= \sum_{(x,y) \in W} [I_x u + I_y v]^2 = \sum_{(x,y) \in W} I_x^2 u^2 + 2I_x I_y u v + I_y^2 v^2 \end{aligned}$$

# Corner Detection: Mathematics

The quadratic approximation can be written as

$$\sum_{(x,y) \in W} I_x^2 u^2 + 2I_x I_y u v + I_y^2 v^2$$

$$E(u, v) \approx [u \quad v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a second moment matrix computed from image derivatives:

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$

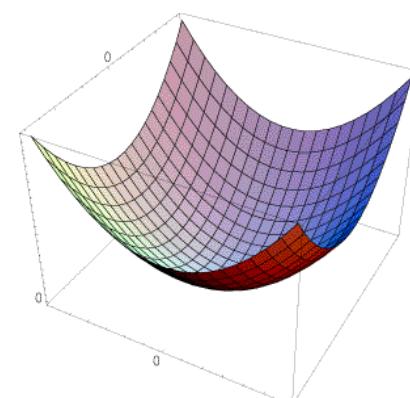
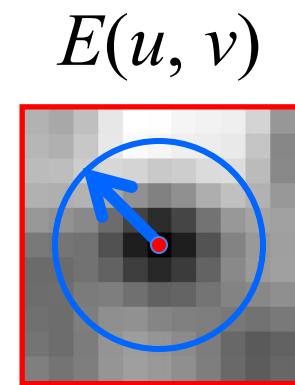
(the sums are over all the pixels in the window  $W$ )

# Interpreting the second moment matrix

- The surface  $E(u,v)$  is locally approximated by a quadratic form. Let's try to understand its shape.
  - Specifically, in which directions does it have the smallest/greatest change?

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$

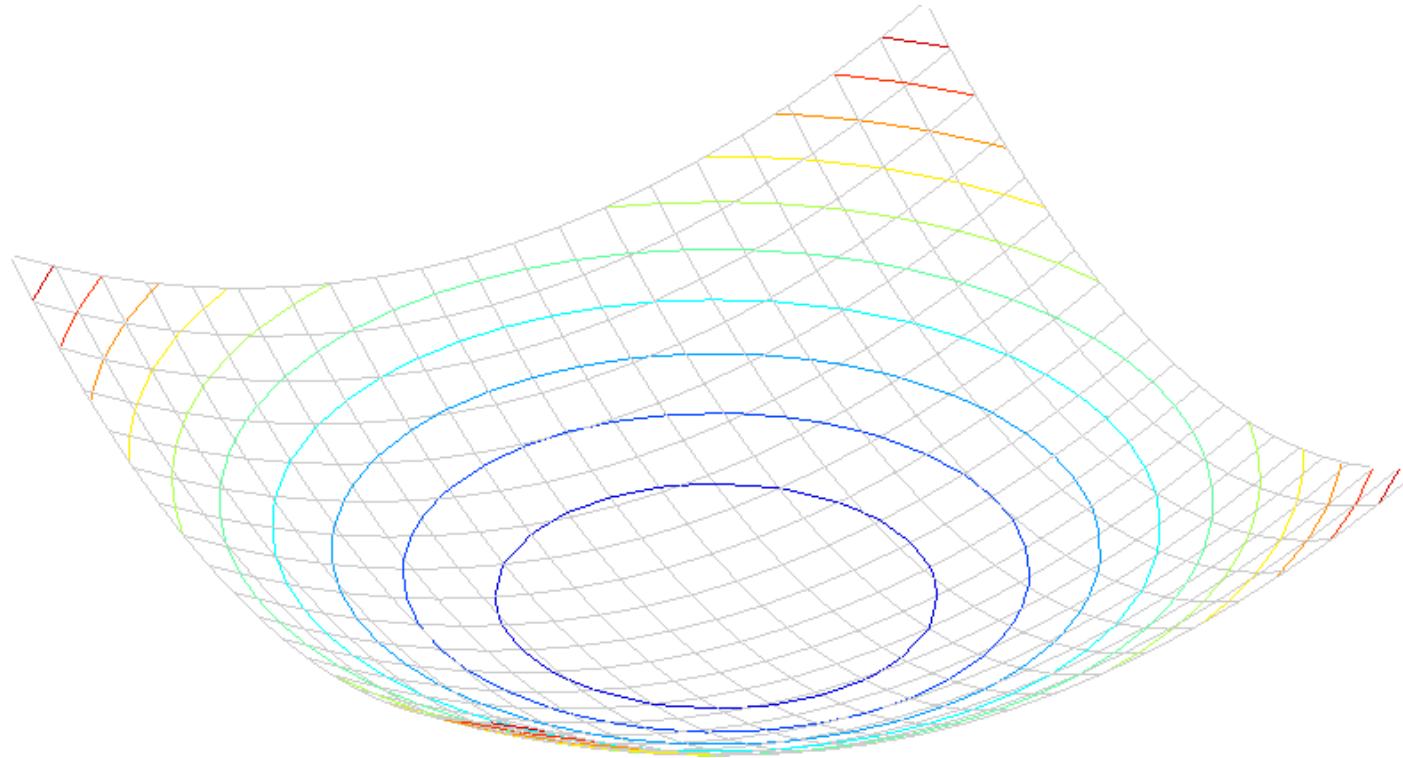


# Interpreting the second moment matrix

Consider a horizontal “slice” of  $E(u, v)$ :

This is the equation of an ellipse.

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$



# Interpreting the second moment matrix

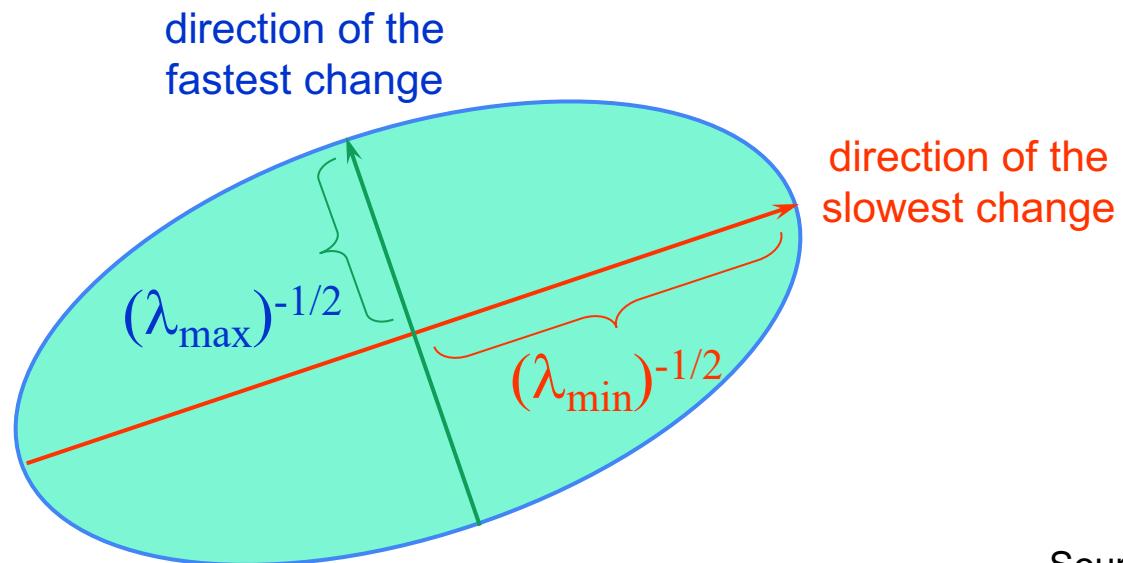
Consider a horizontal “slice” of  $E(u, v)$ :

This is the equation of an ellipse.

Diagonalization of  $M$ :

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$
$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$



## Interpreting the second moment matrix

Consider the axis-aligned case (gradients are either horizontal or vertical)

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

If either  $a$  or  $b$  is close to 0, then this is not a corner, so look for locations where both are large.

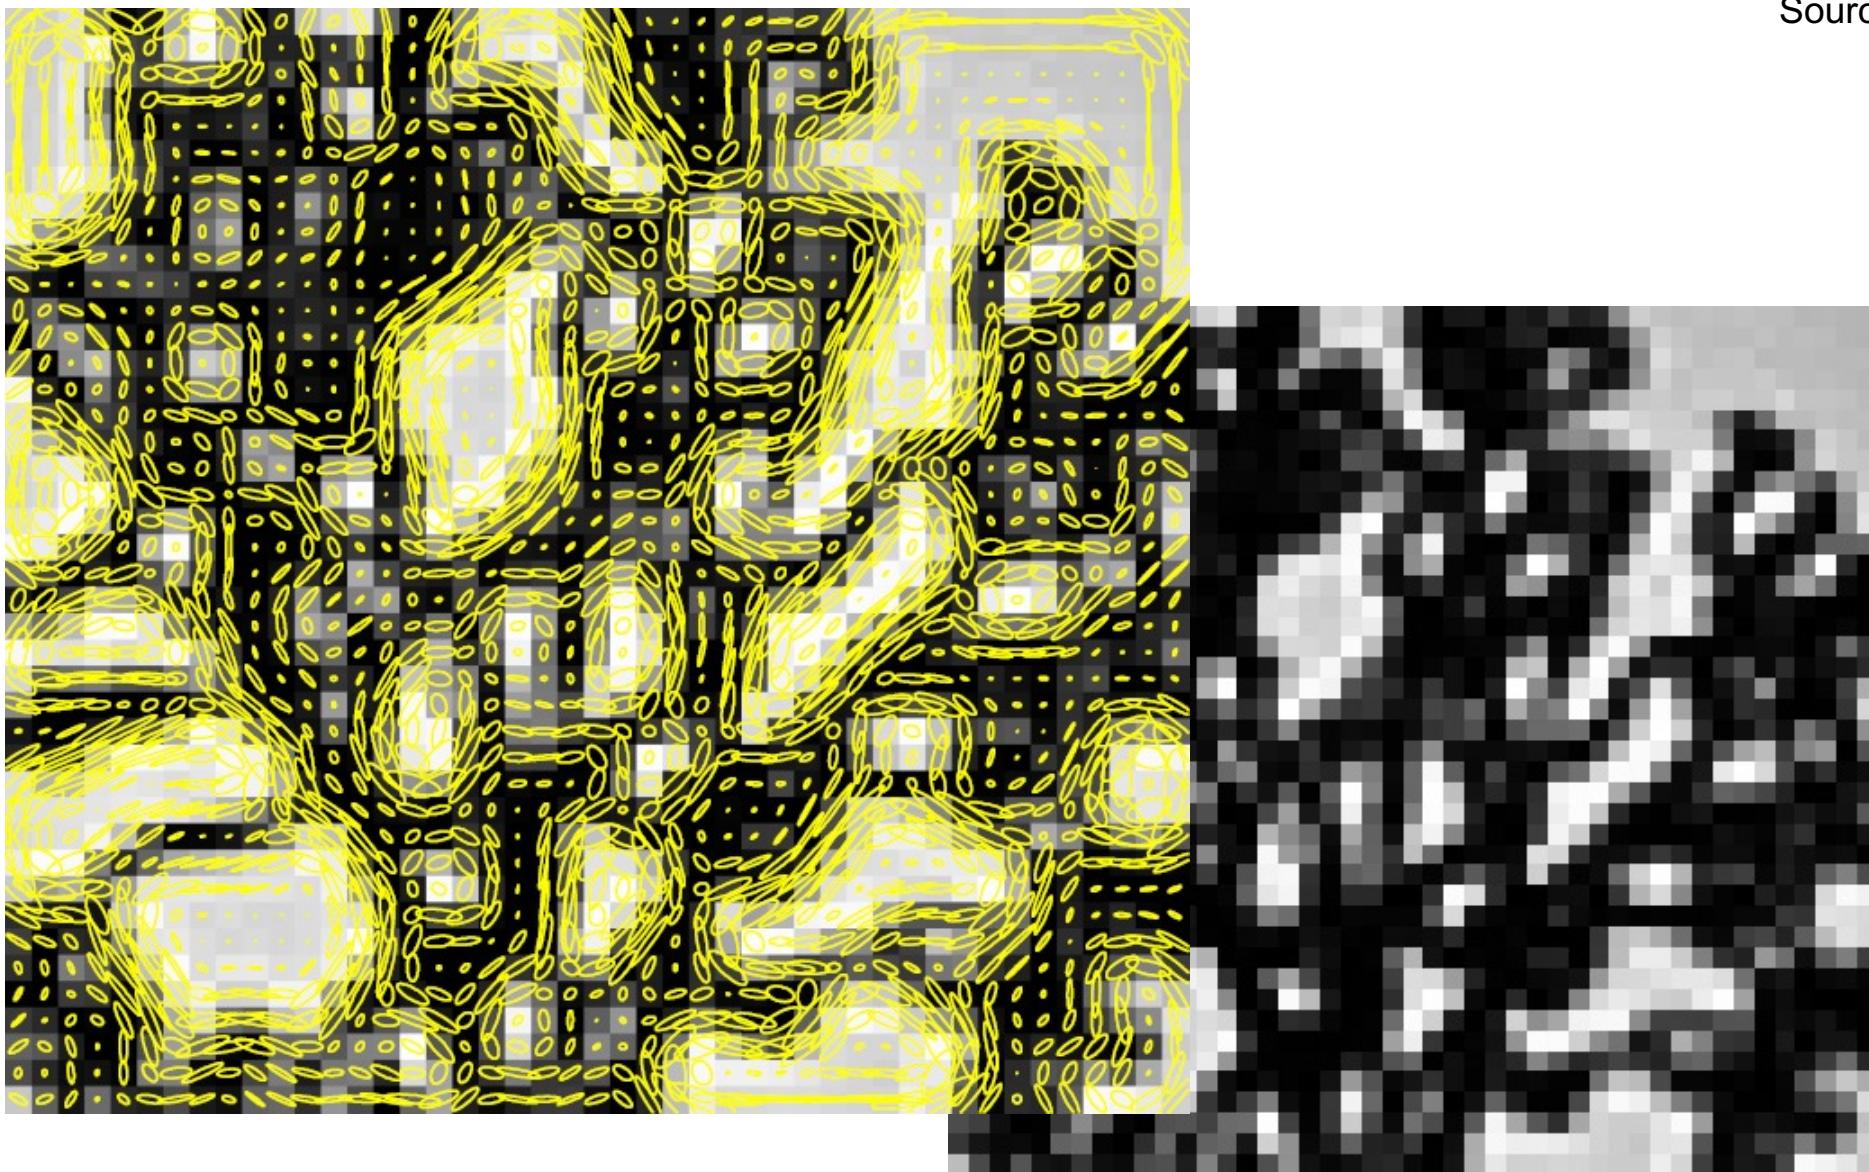
# Visualization of second moment matrices

Source: S. Lazebnik



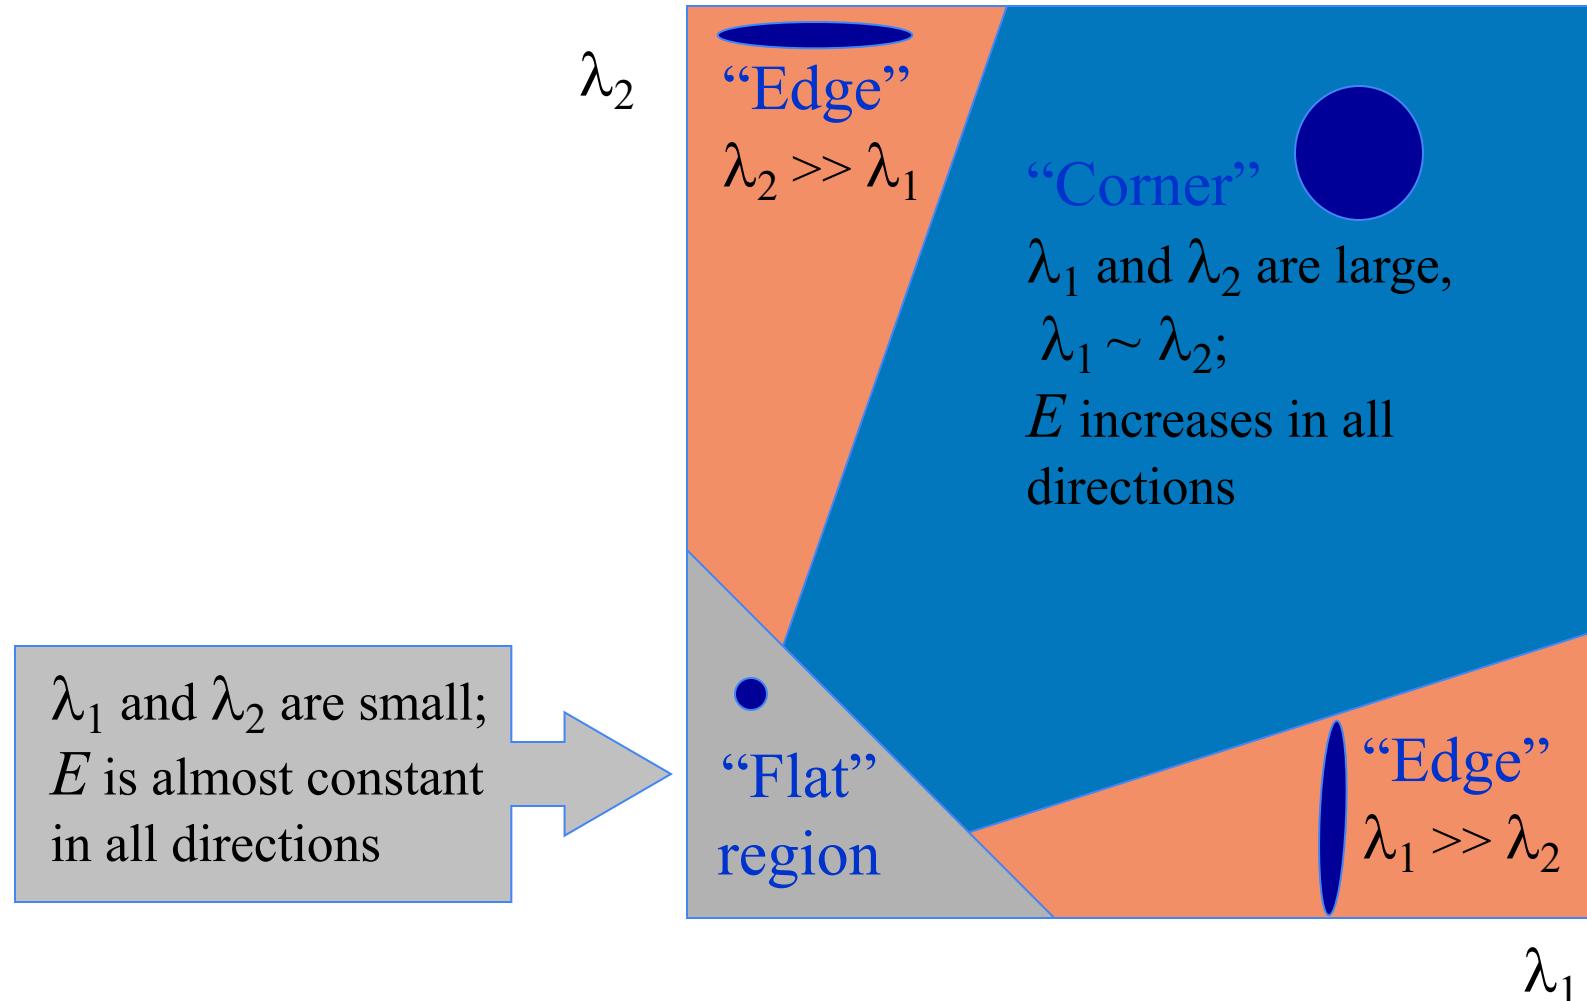
# Visualization of second moment matrices

Source: S. Lazebnik



# Interpreting the eigenvalues

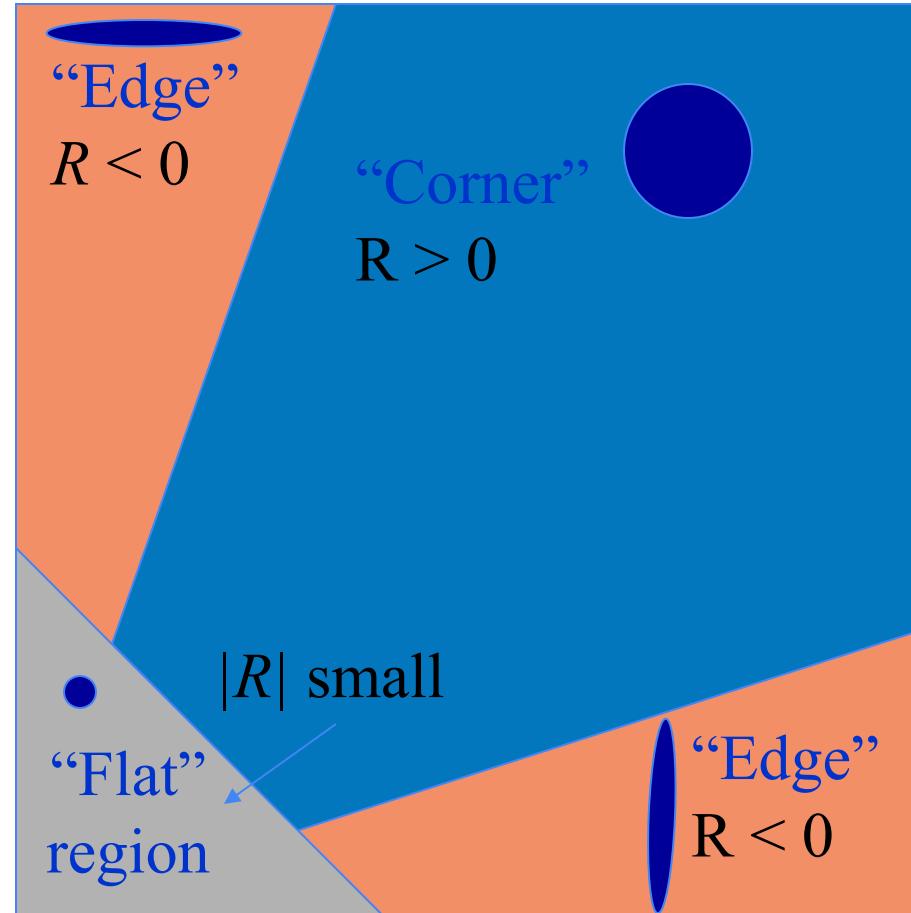
Classification of image points using eigenvalues of  $M$ :



# Corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : constant (0.04 to 0.06)



# Keypoints – Harris corner detector

# The Harris corner detector

- Compute partial derivatives at each pixel
- Compute second moment matrix M in a Gaussian window around each pixel:

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens. [A Combined Corner and Edge Detector.](#)  
*Proceedings of the 4th Alvey Vision Conference:* pages 147—151, 1988.

# The Harris corner detector

- Compute partial derivatives at each pixel
- Compute second moment matrix M in a Gaussian window around each pixel
- Compute corner response function R

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)  
*Proceedings of the 4th Alvey Vision Conference:* pages 147—151, 1988.

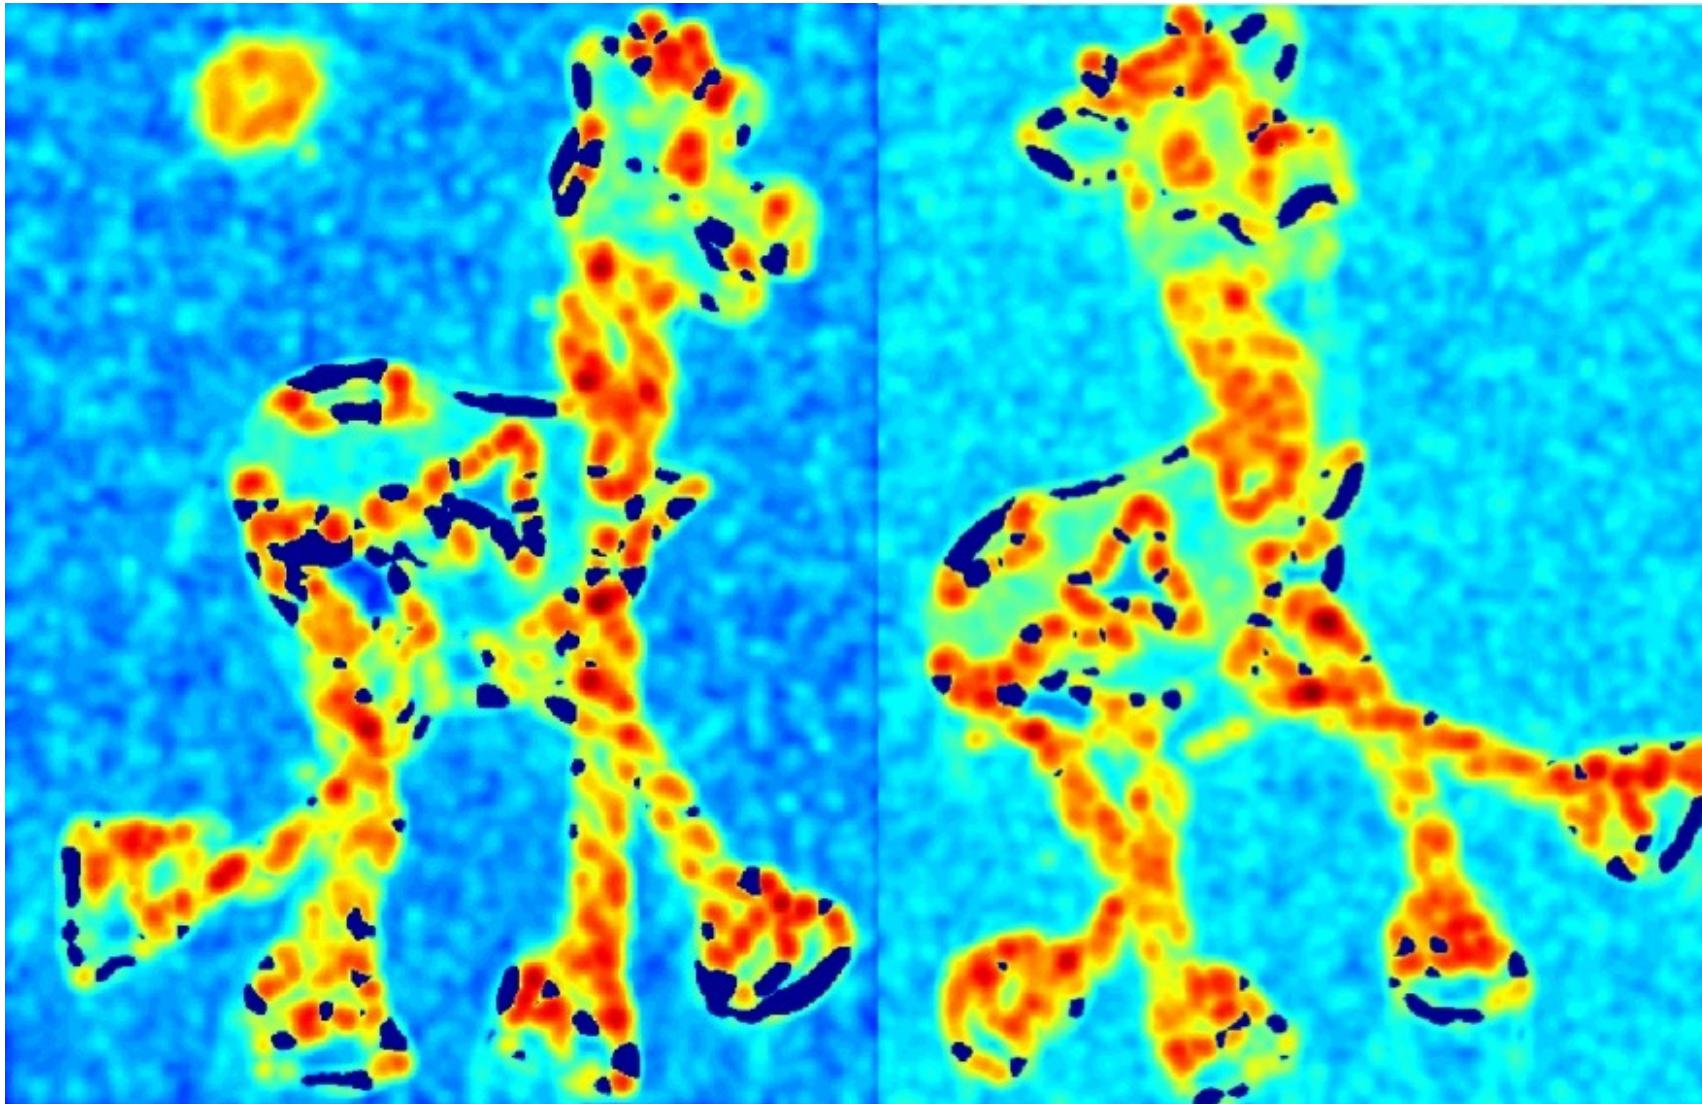
# Harris Detector: Steps



Source: S. Lazebnik

# Harris Detector: Steps

Compute corner response R



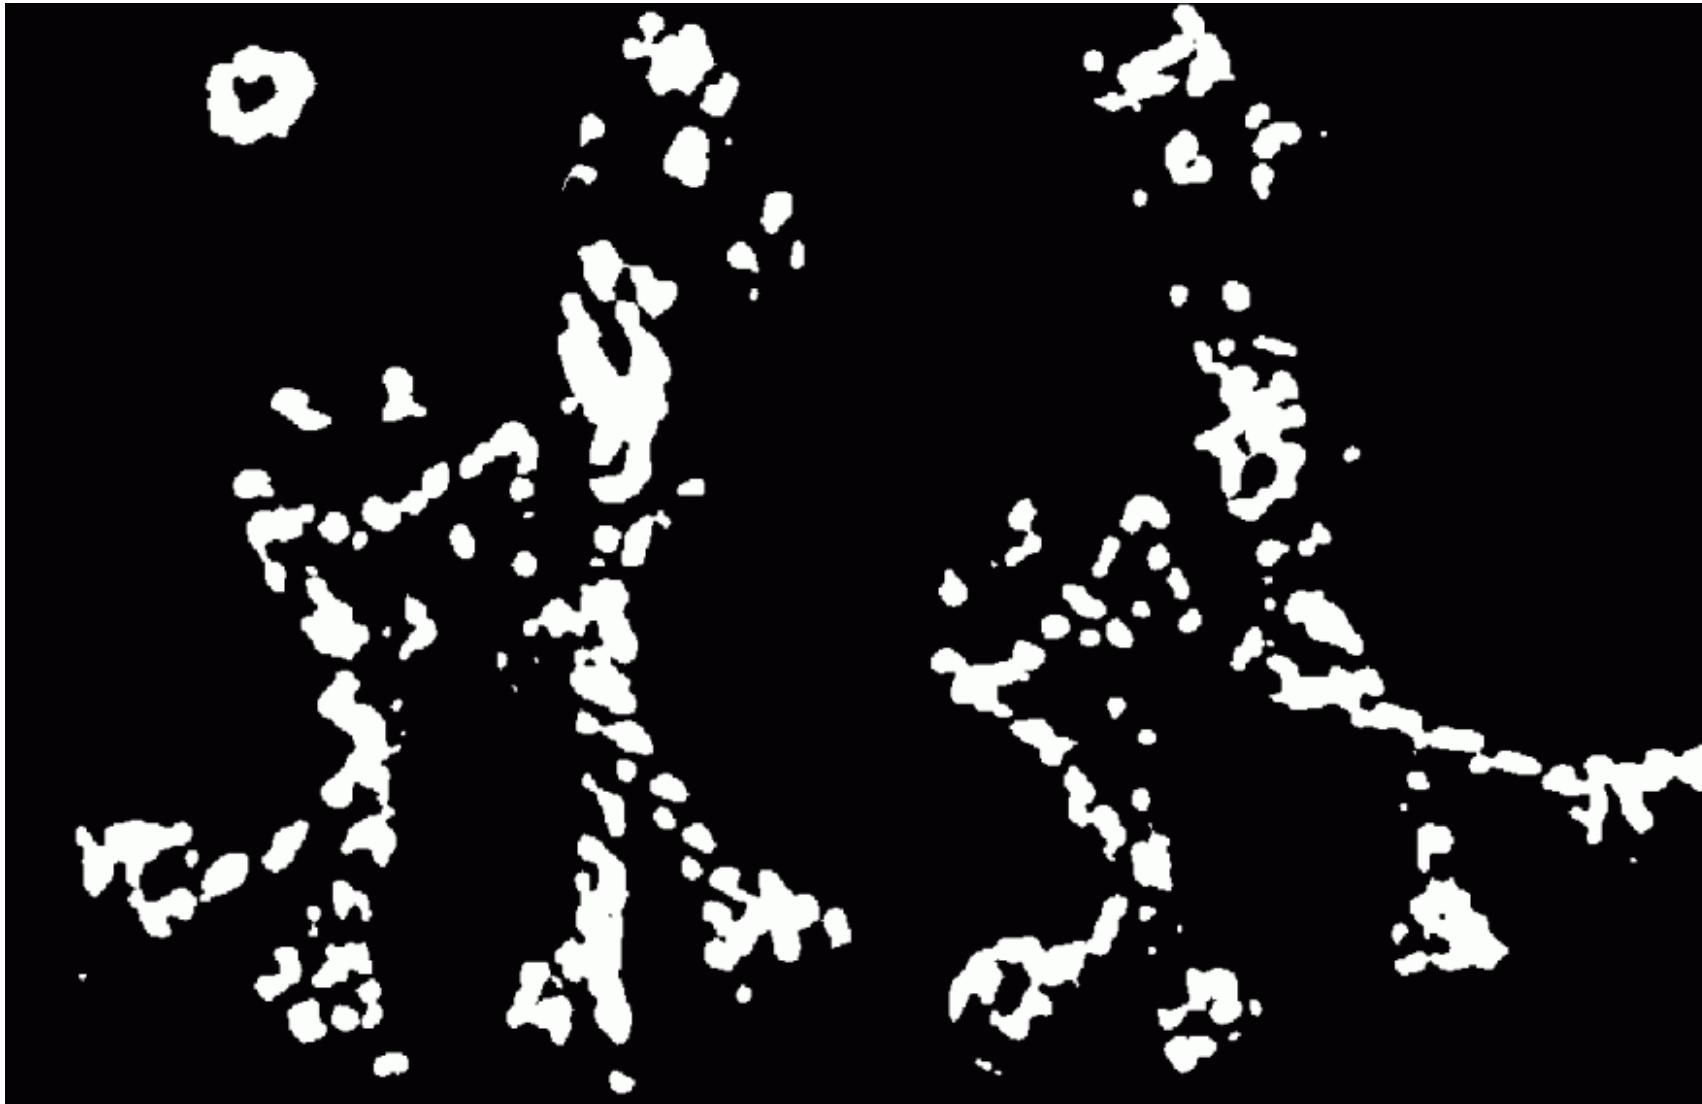
# The Harris corner detector

- Compute partial derivatives at each pixel
- Compute second moment matrix  $M$  in a Gaussian window around each pixel
- Compute corner response function  $R$
- Threshold  $R$
- Find local maxima of response function (nonmaximum suppression)

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)  
*Proceedings of the 4th Alvey Vision Conference:* pages 147—151, 1988.

# Harris Detector: Steps

Find points with large corner response:  $R > \text{threshold}$



# Harris Detector: Steps

Take only the points of local maxima of R



# Harris Detector: Steps



Source: S. Lazebnik

# Robustness of corner features

- What happens to corner features when the image undergoes geometric or photometric transformations?

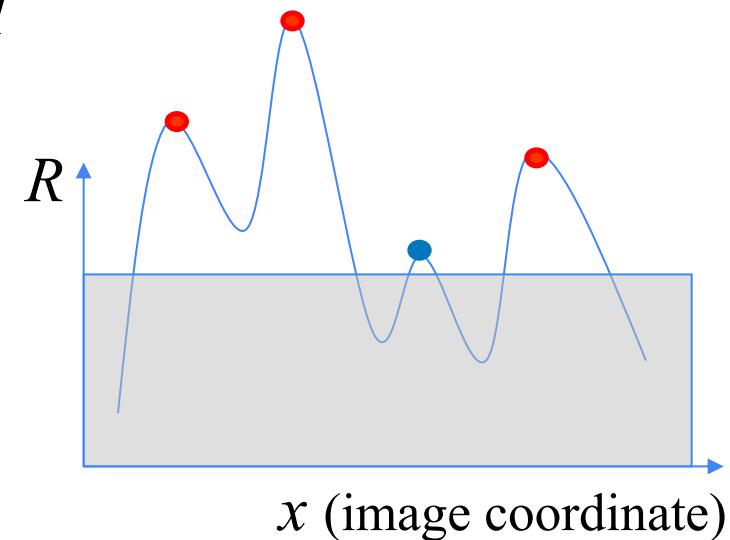
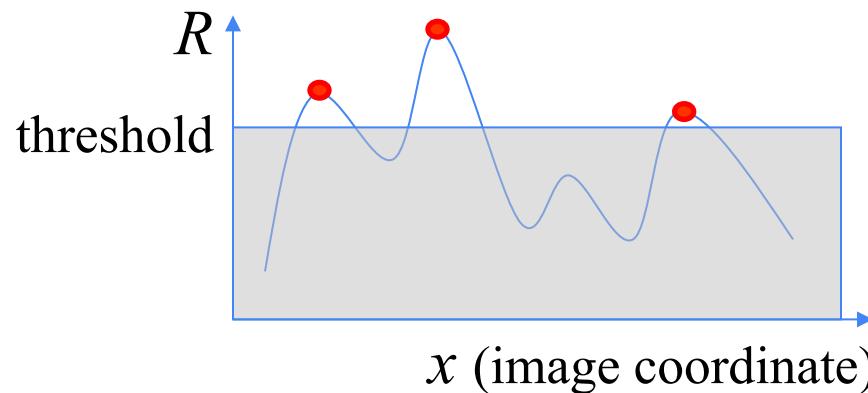


# Affine intensity change



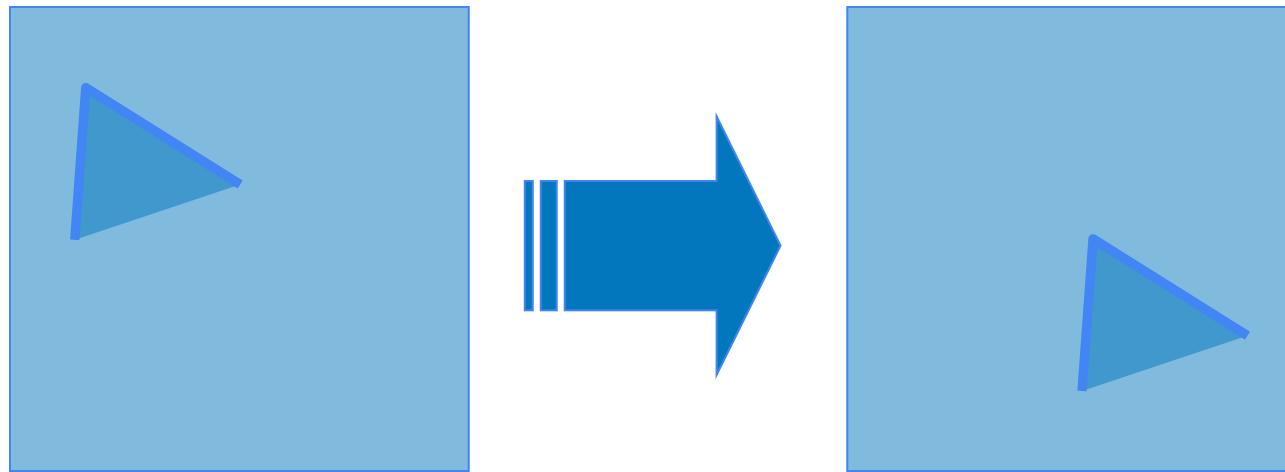
$$I \rightarrow a I + b$$

- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

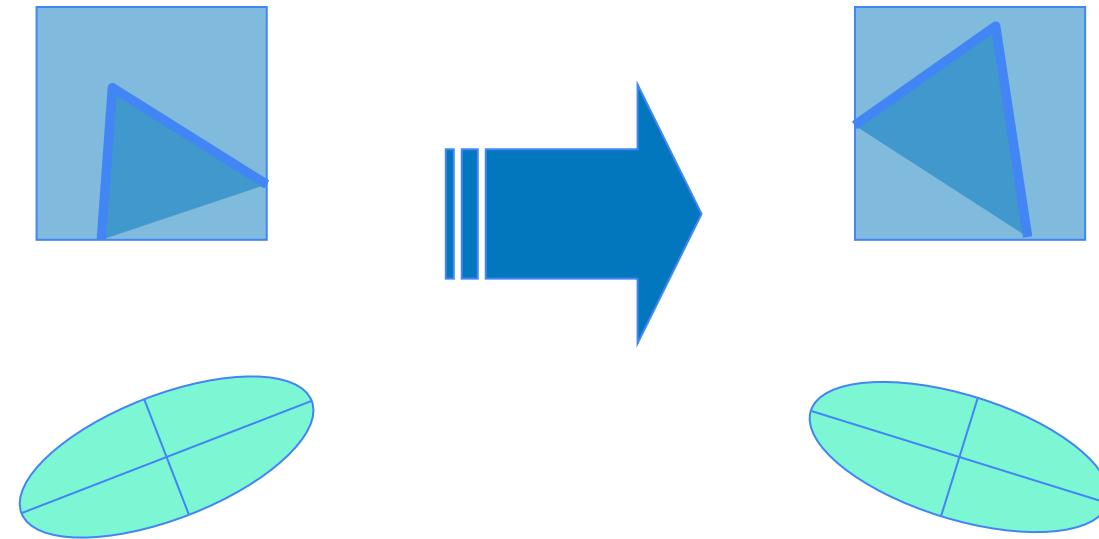
# Image translation



- Derivatives and window function are shift-invariant

Corner location is *covariant* w.r.t. translation

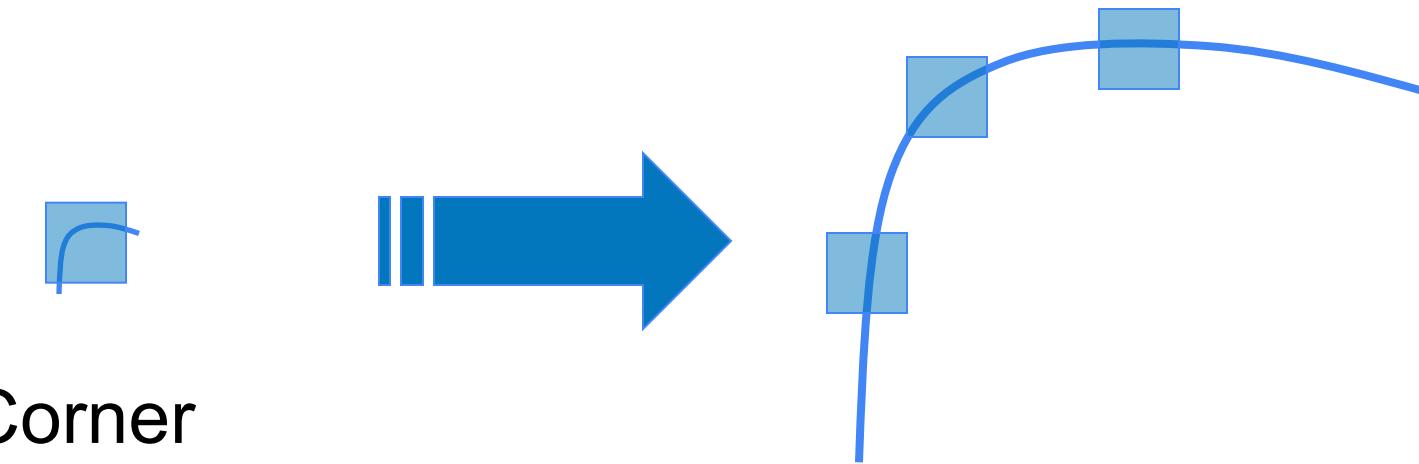
# Image rotation



Second moment ellipse rotates but its shape  
(i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

# Scaling

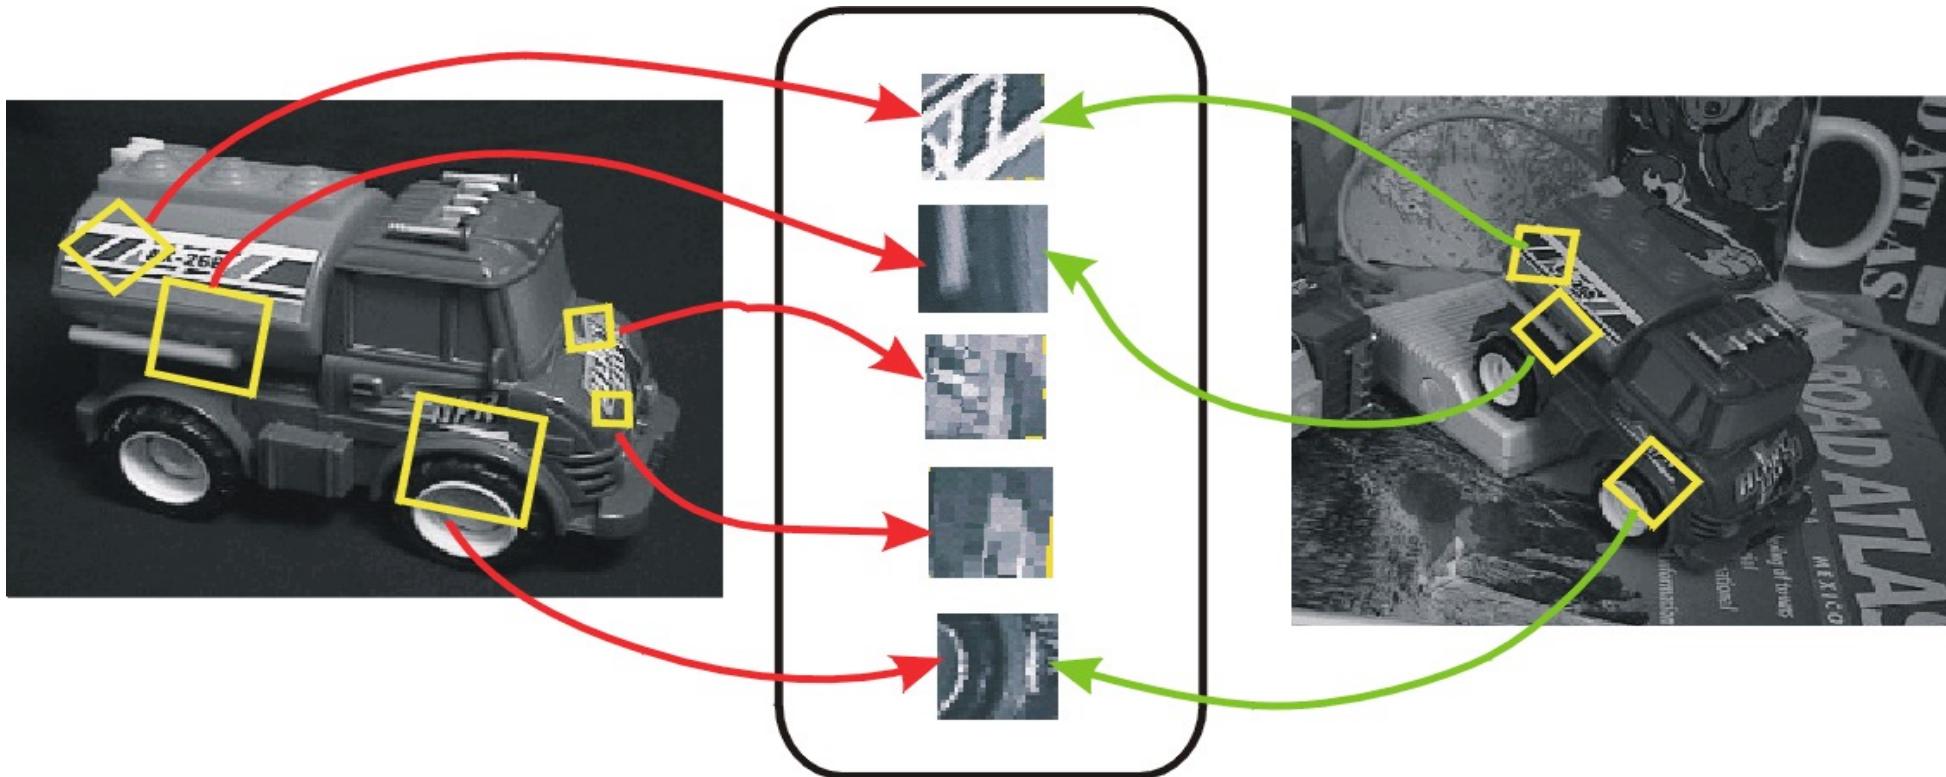


All points will  
be classified  
as **edges**

Corner location is not covariant to scaling!

# Sneak peak: from feature detection to feature description

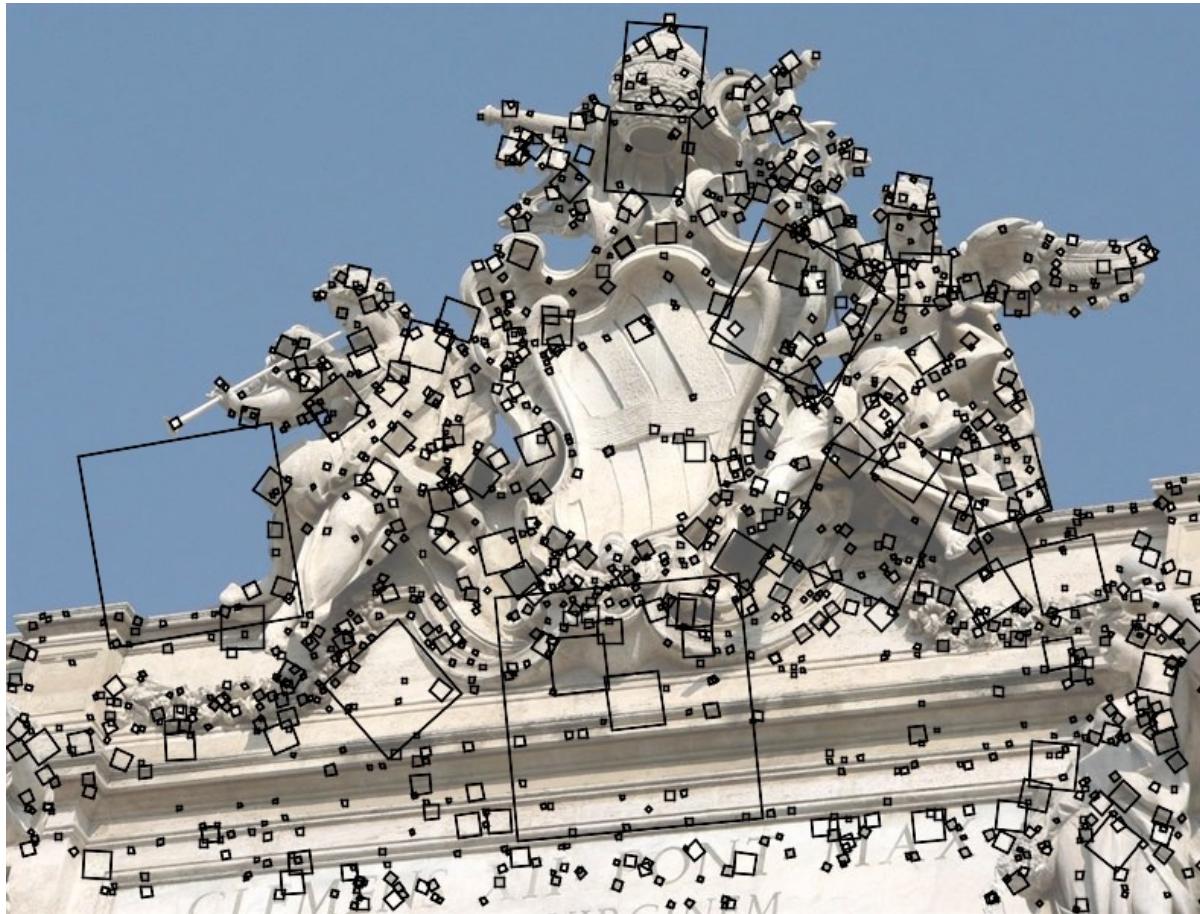
- Detection is covariant:
  - $\text{features}(\text{transform}(\text{image})) = \text{transform}(\text{features}(\text{image}))$



# Keypoints – Scale Invariant Feature Transform (SIFT)



# SIFT keypoint detection

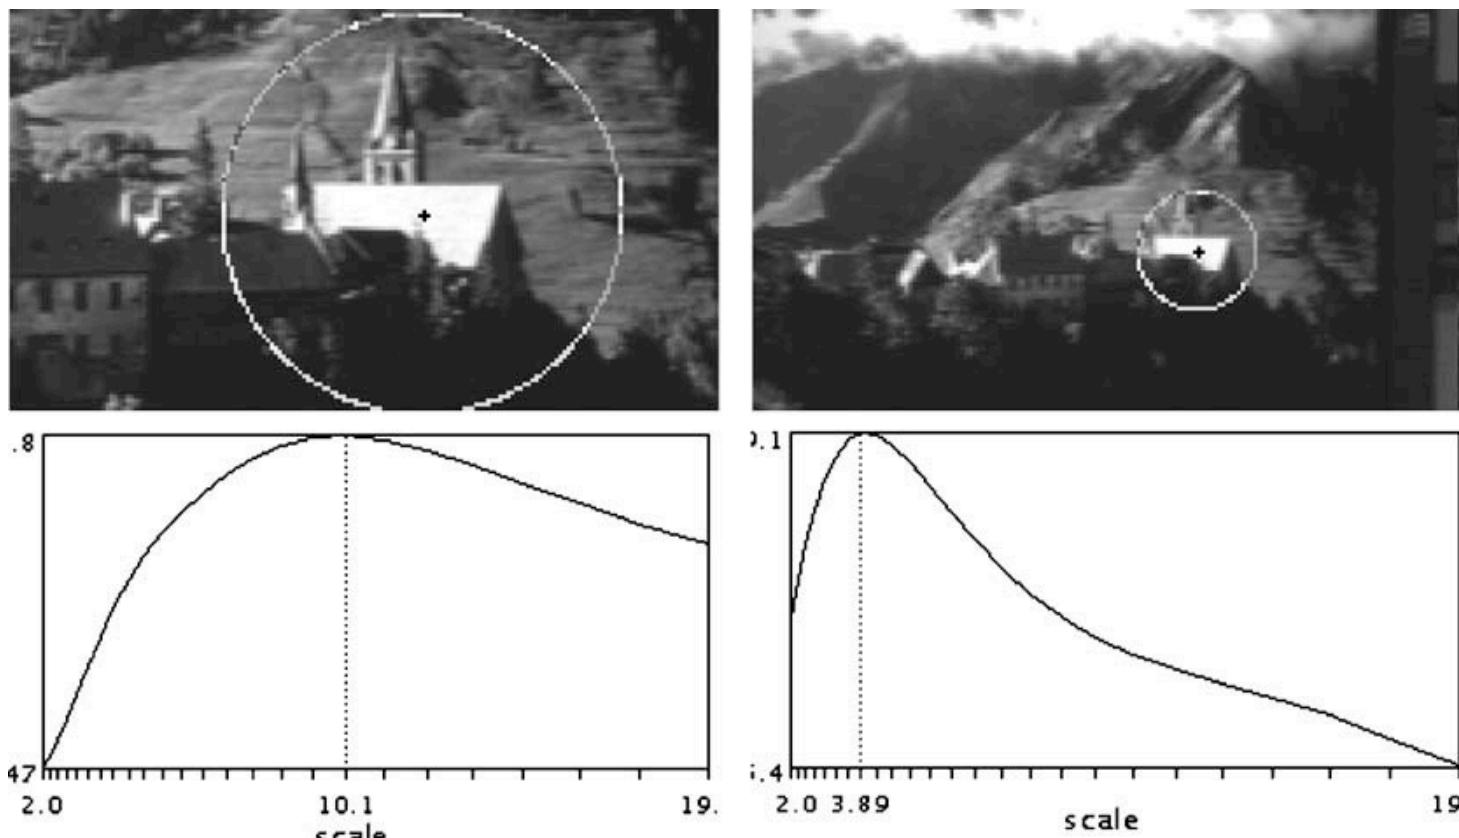


D. Lowe, [\*\*Distinctive image features from scale-invariant keypoints\*\*](#),  
*IJCV* 60 (2), pp. 91-110, 2004.

Source: S. Lazebnik

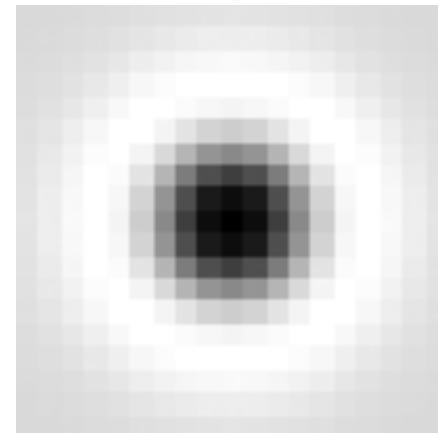
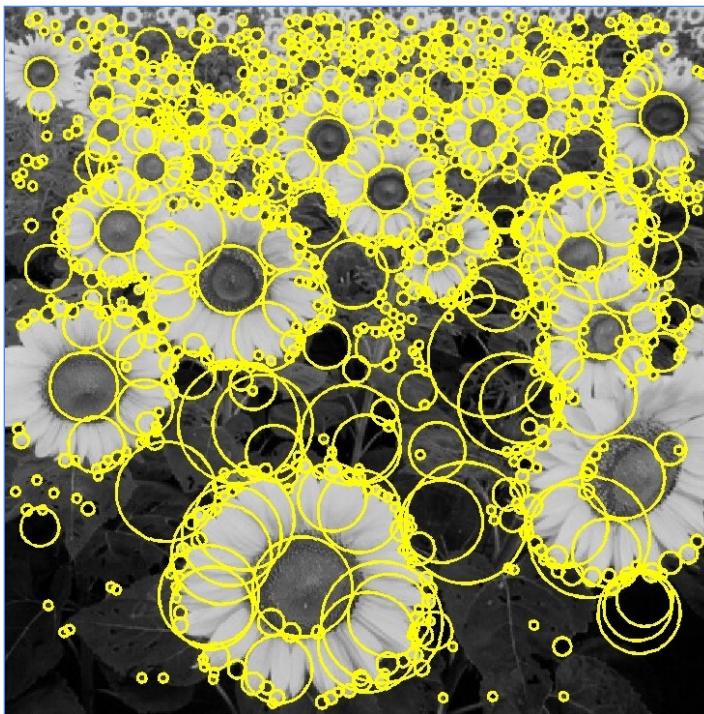
# Keypoint detection with scale selection

- We want to extract keypoints with characteristic scale that is covariant with the image transformation



# Basic idea

- Convolve the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting scale space



T. Lindeberg. [Feature detection with automatic scale selection.](#)  
*IJCV* 30(2), pp 77-116, 1998.

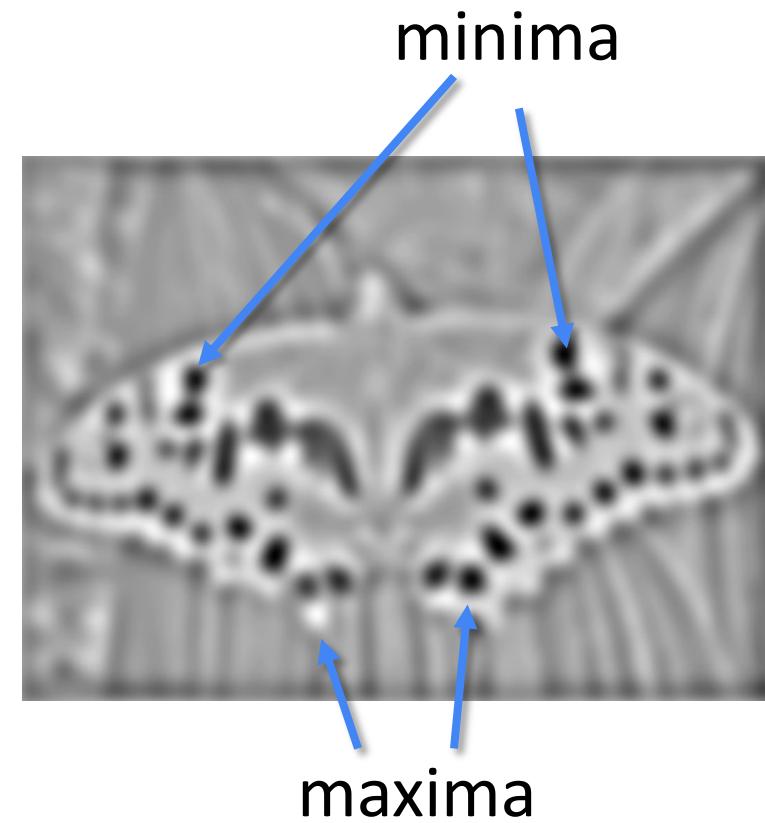
Source: S. Lazebnik

# Blob detection

- Find maxima and minima of blob filter response in space and scale

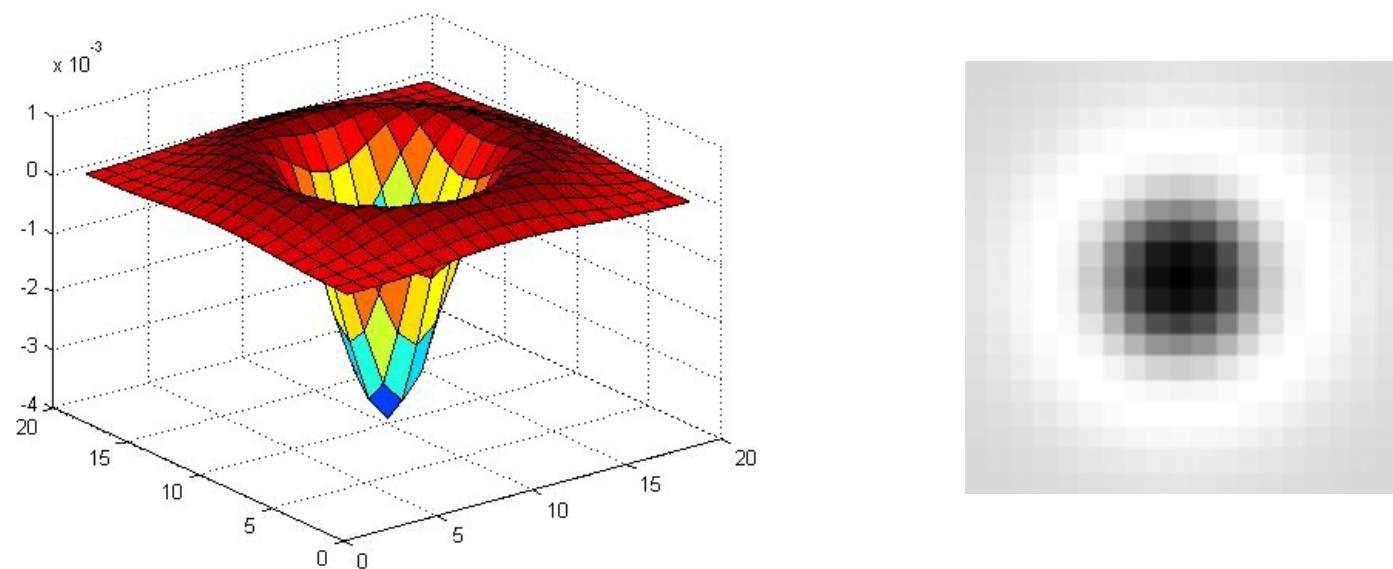


$$* \quad \bullet =$$



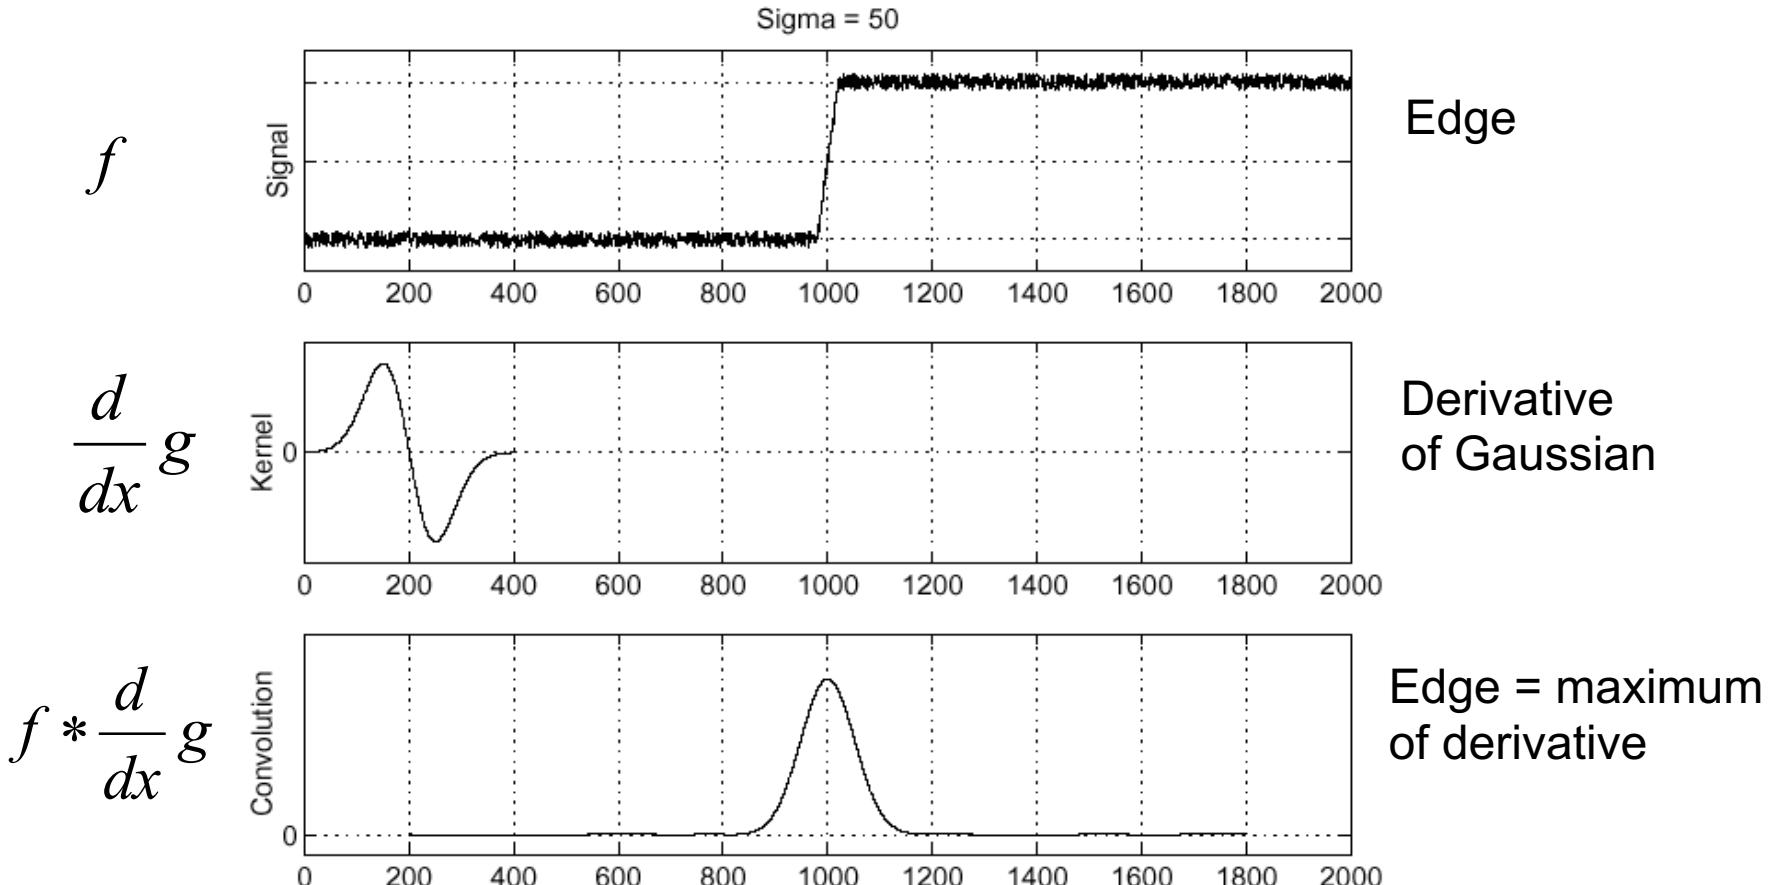
# Blob filter

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

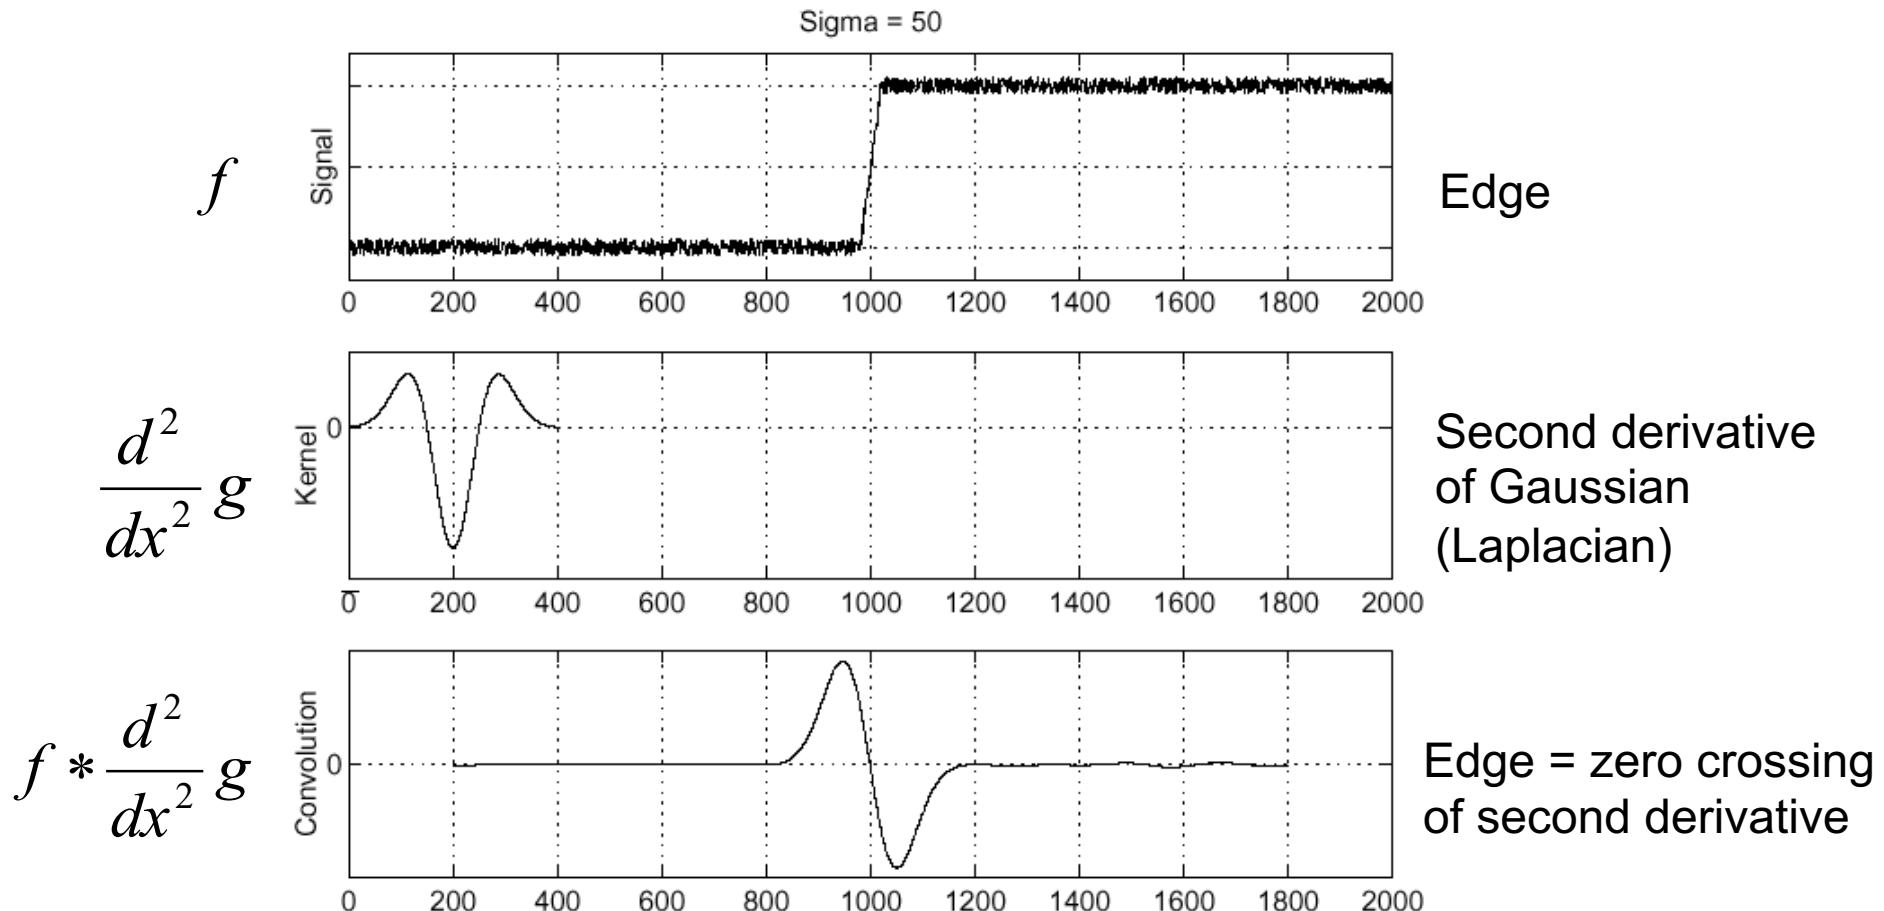


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Recall: Edge detection

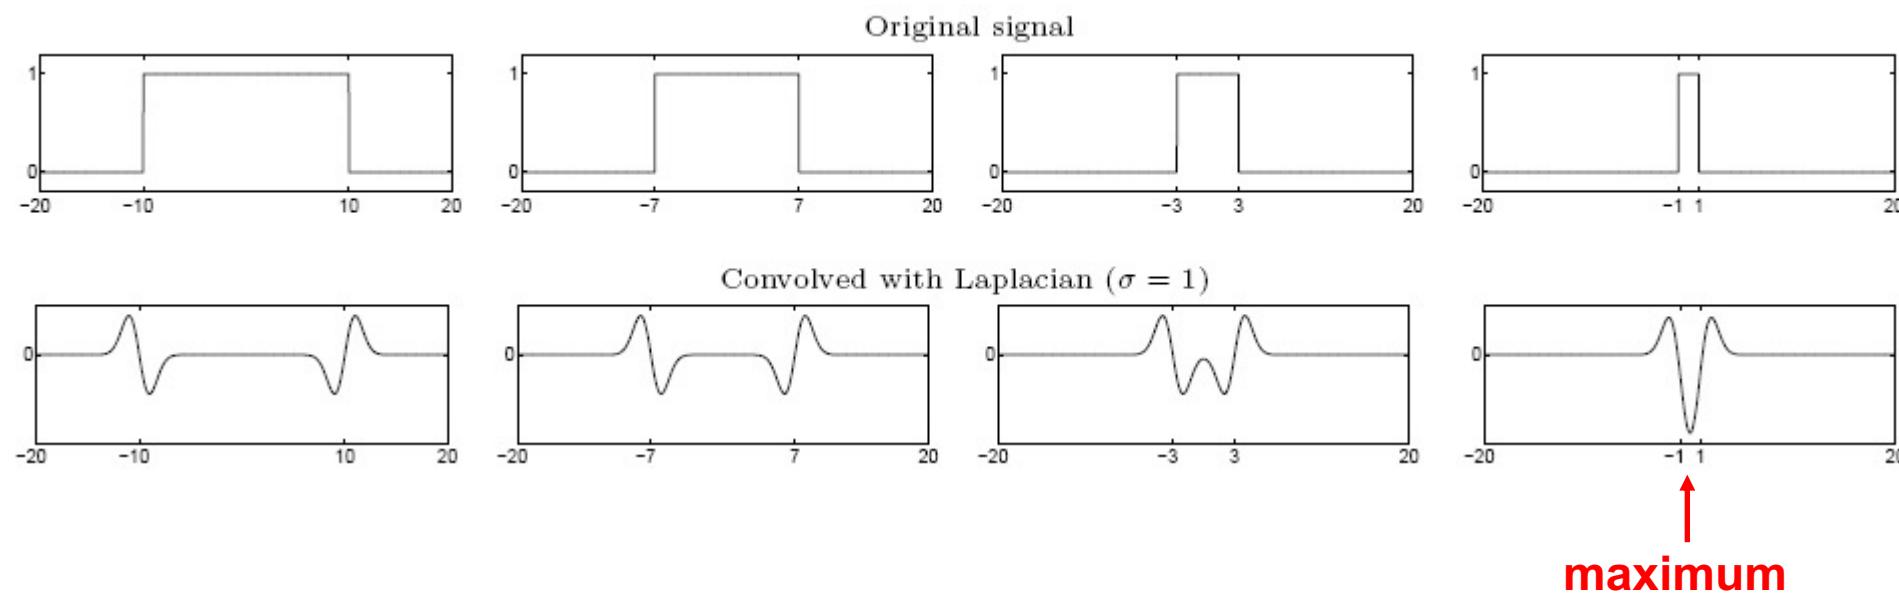


# Edge detection, Take 2



# From edges to blobs

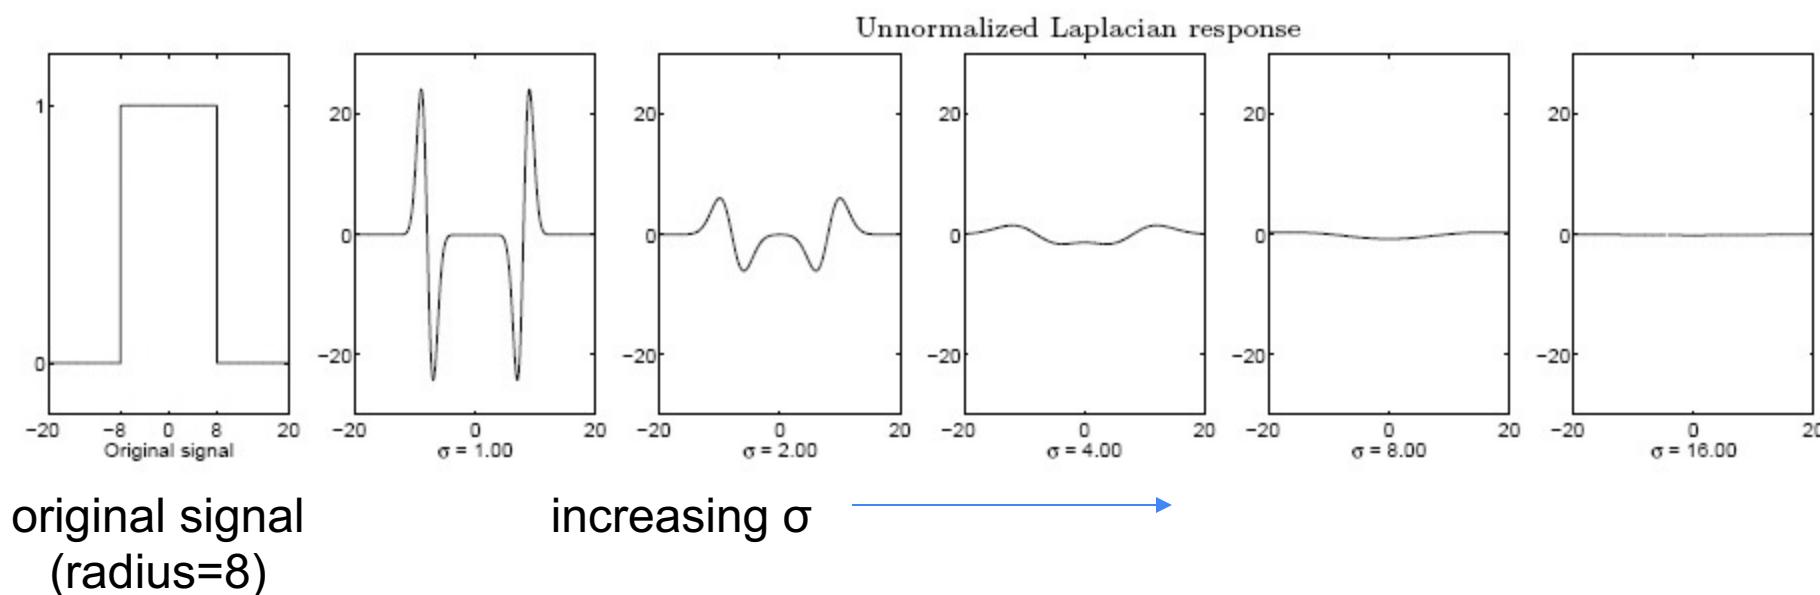
- Edge = ripple
- Blob = superposition of two ripples



**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

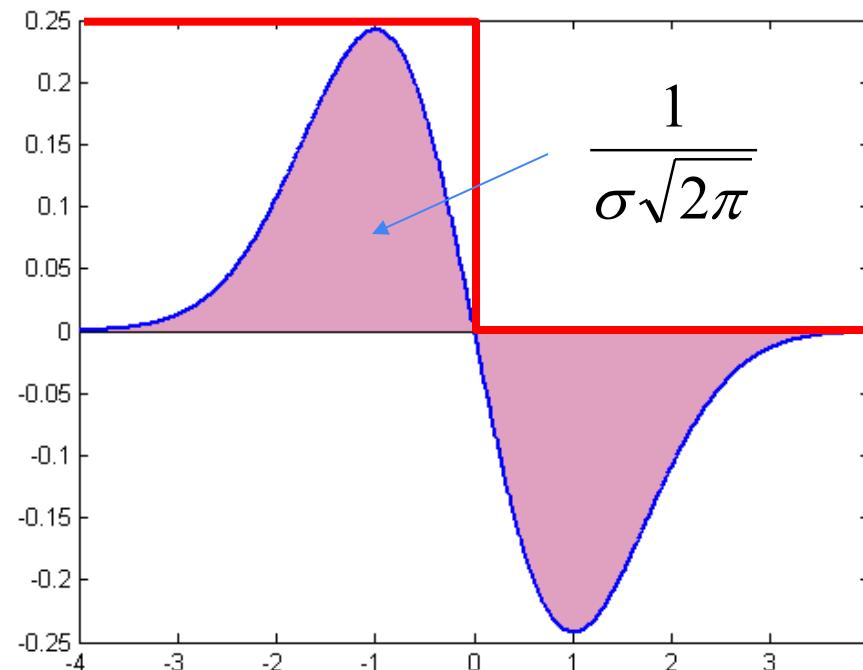
# Scale selection

- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:



# Scale normalization

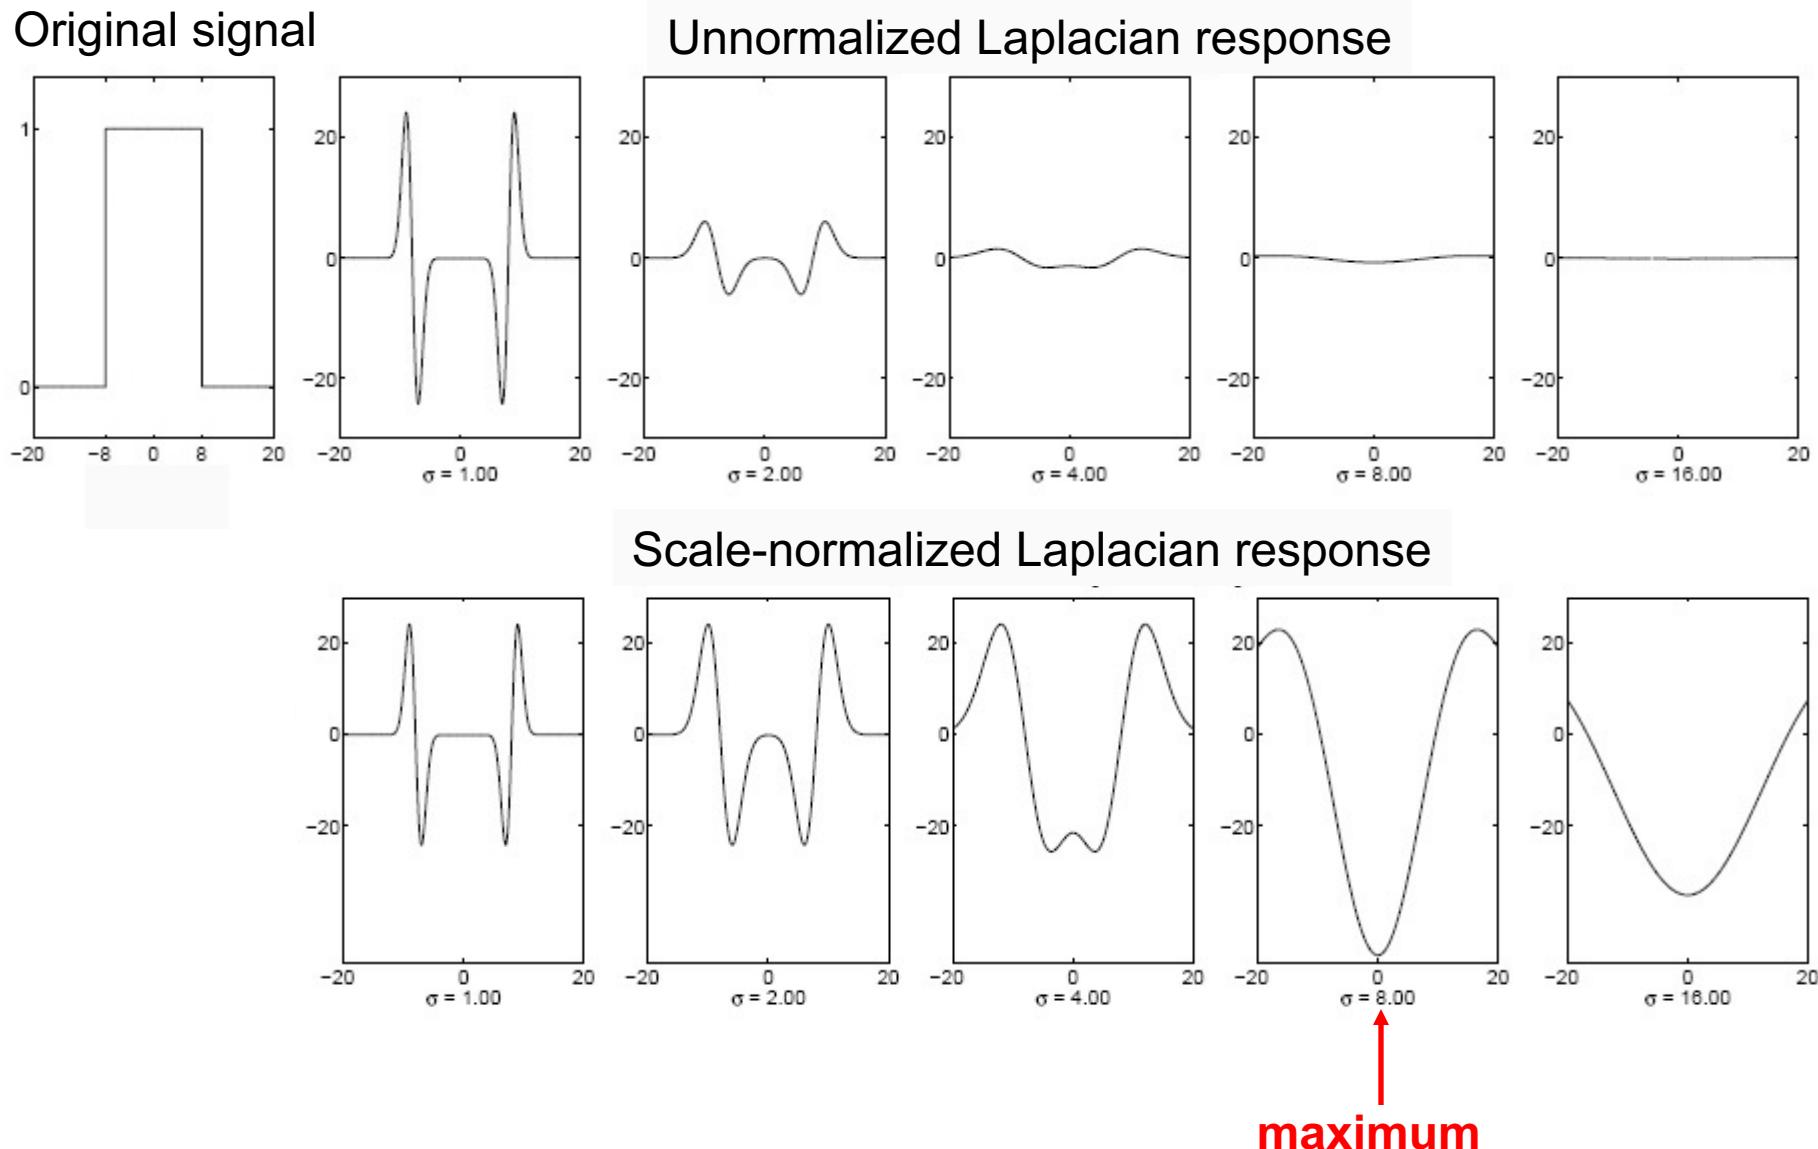
- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases



# Scale normalization

- The response of a derivative of Gaussian filter to a perfect step edge decreases as  $\sigma$  increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by  $\sigma$
- Laplacian is the second Gaussian derivative, so it must be multiplied by  $\sigma^2$

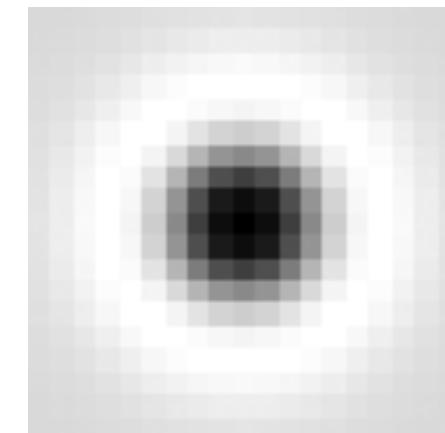
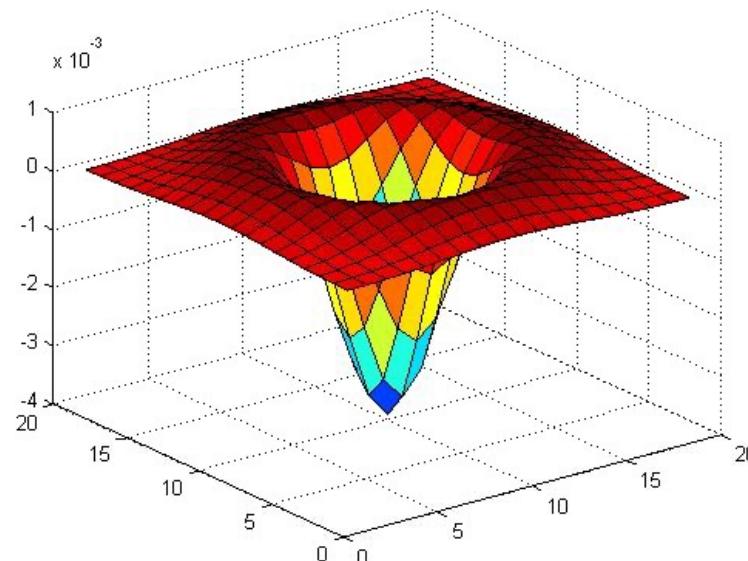
# Effect of scale normalization



# Blob detection in 2D

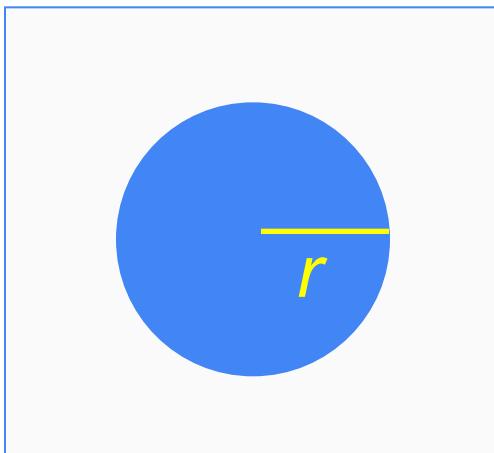
- Scale-normalized Laplacian of Gaussian:

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left( \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

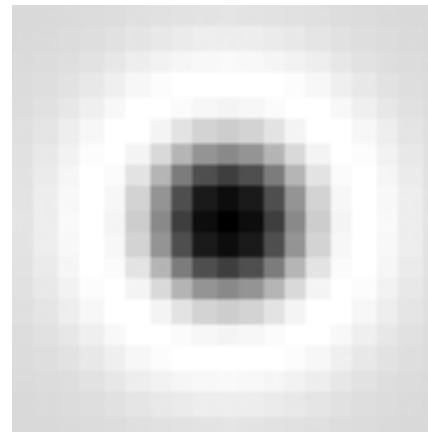


# Blob detection in 2D

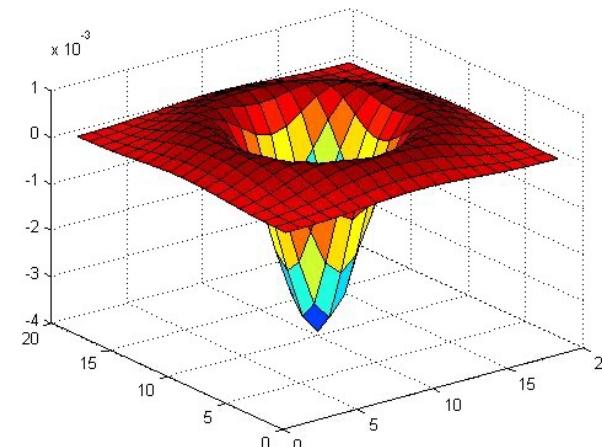
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?



image



Laplacian



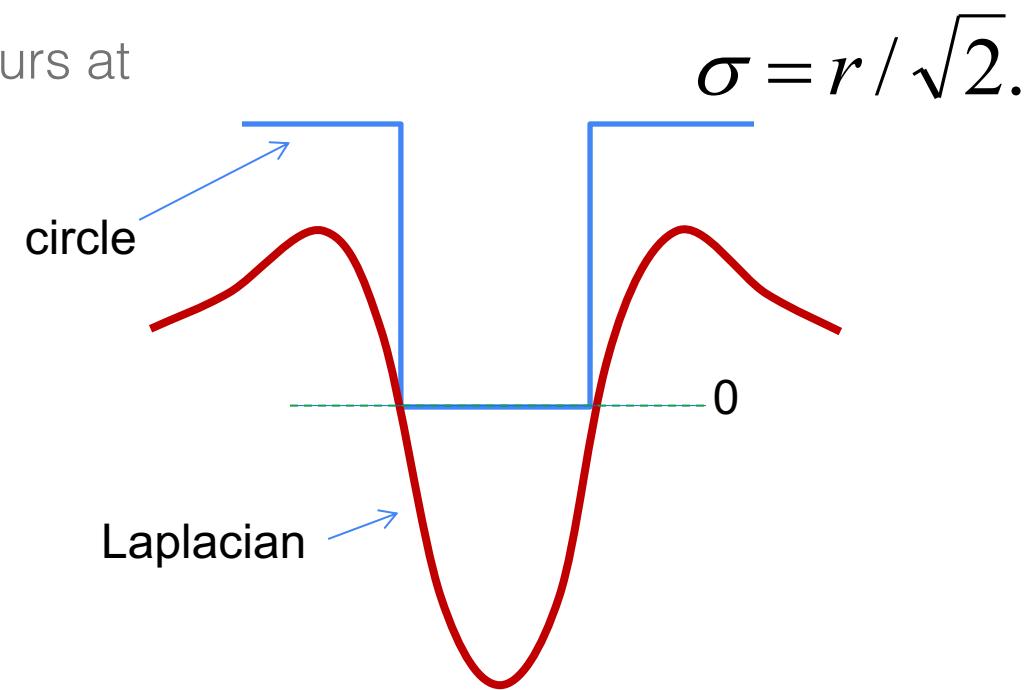
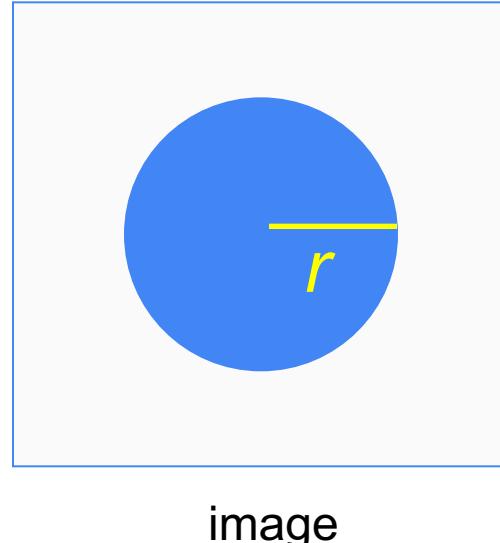
Source: S. Lazebnik

# Blob detection in 2D

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius  $r$ ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$

- Therefore, the maximum response occurs at



# Scale-space blob detector

- Convolve image with scale-normalized Laplacian at several scales

# Scale-space blob detector: Example



Source: S. Lazebnik

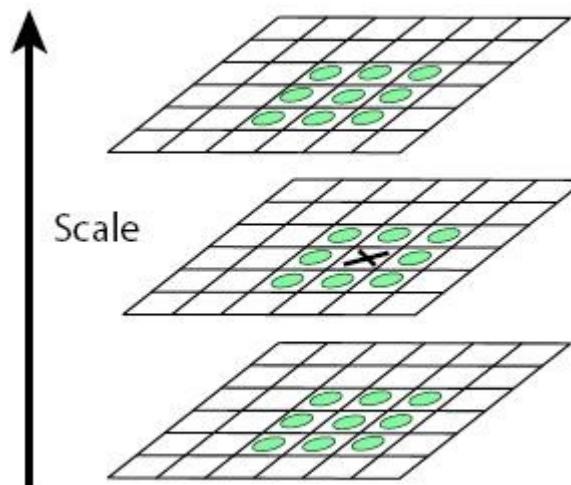
# Scale-space blob detector: Example



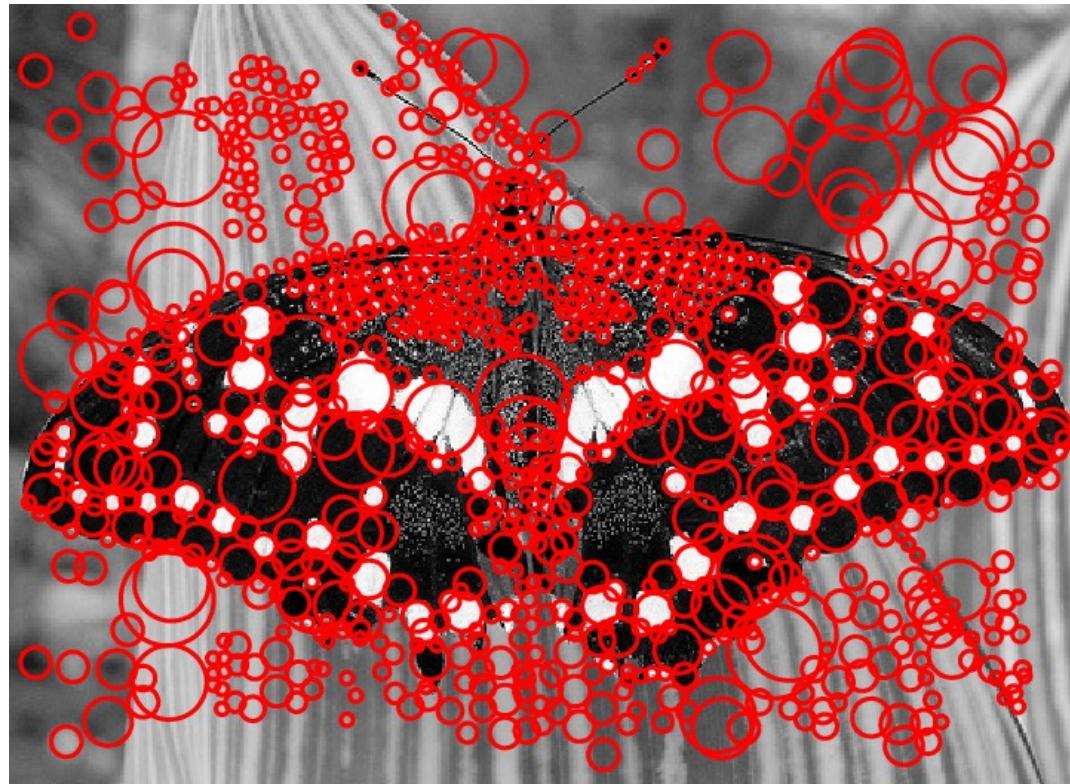
sigma = 11.9912

# Scale-space blob detector

- Convolve image with scale-normalized Laplacian at several scales
- Find maxima of squared Laplacian response in scale-space



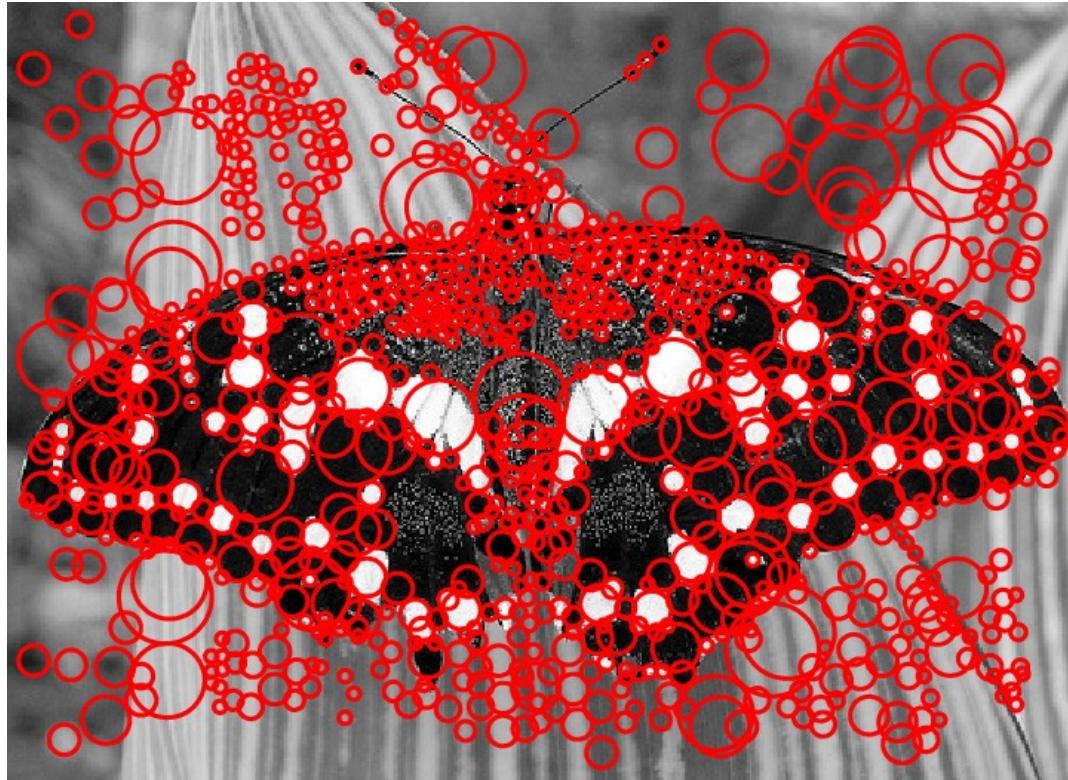
# Scale-space blob detector: Example



Source: S. Lazebnik

# Eliminating edge responses

- Laplacian has strong response along edge



# Eliminating edge responses

- Laplacian has strong response along edge



- Solution: filter based on Harris response function over neighborhoods containing the “blobs”

# Efficient implementation

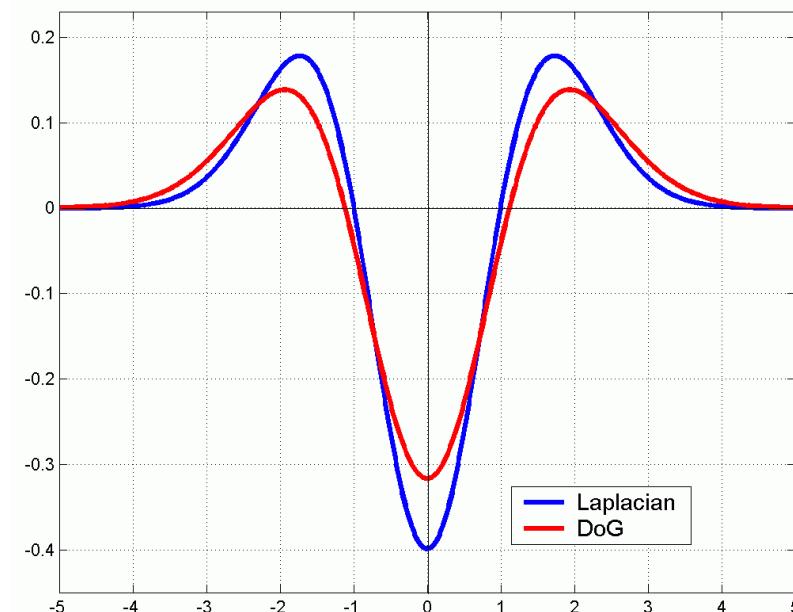
- Approximating the Laplacian with a difference of Gaussians:

$$L = \sigma^2 \left( G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

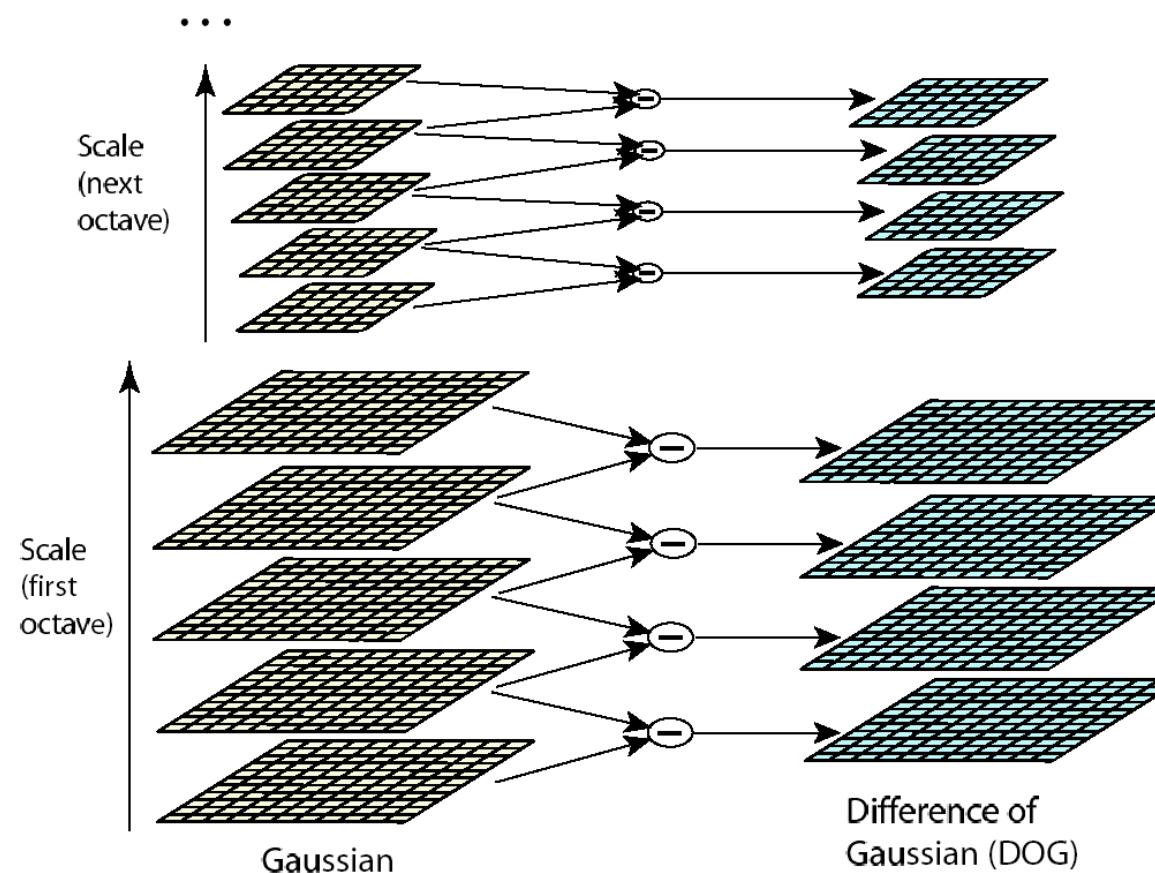
(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



# Efficient implementation



David G. Lowe. [\*\*"Distinctive image features from scale-invariant keypoints."\*\*](#) IJCV 60 (2), pp. 91-110, 2004.

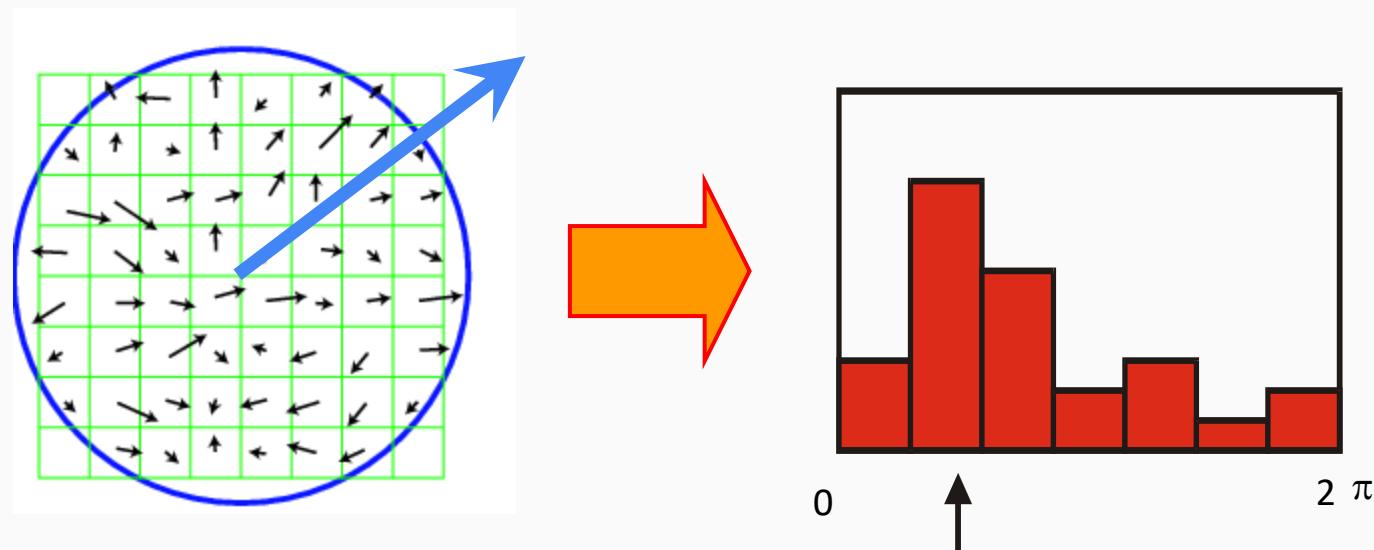
# From feature detection to feature description

- Scaled and rotated versions of the same neighborhood will give rise to blobs that are related by the same transformation
- What to do if we want to compare the appearance of these image regions?
  - Normalization: transform these regions into same-size circles
  - Problem: rotational ambiguity



# Eliminating rotation ambiguity

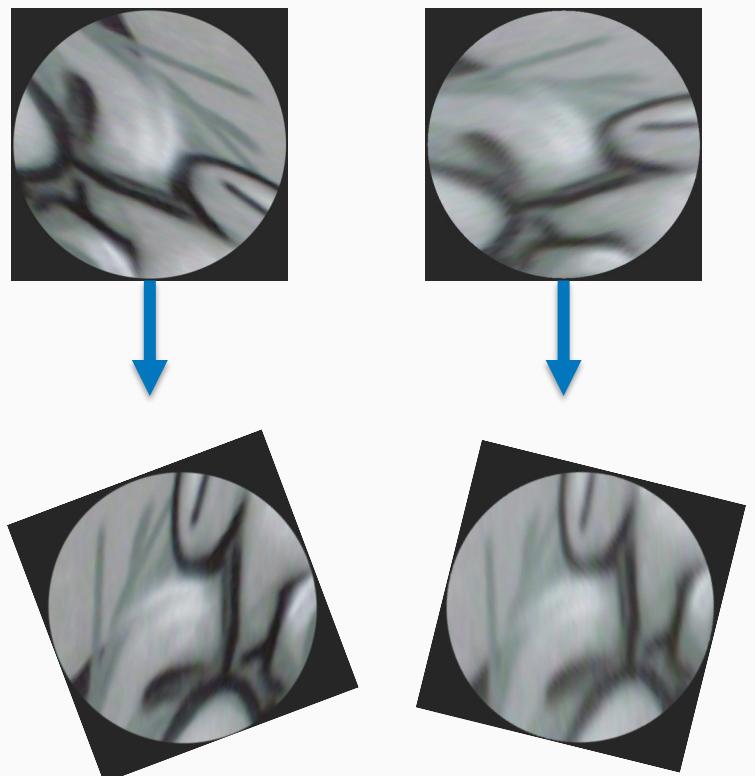
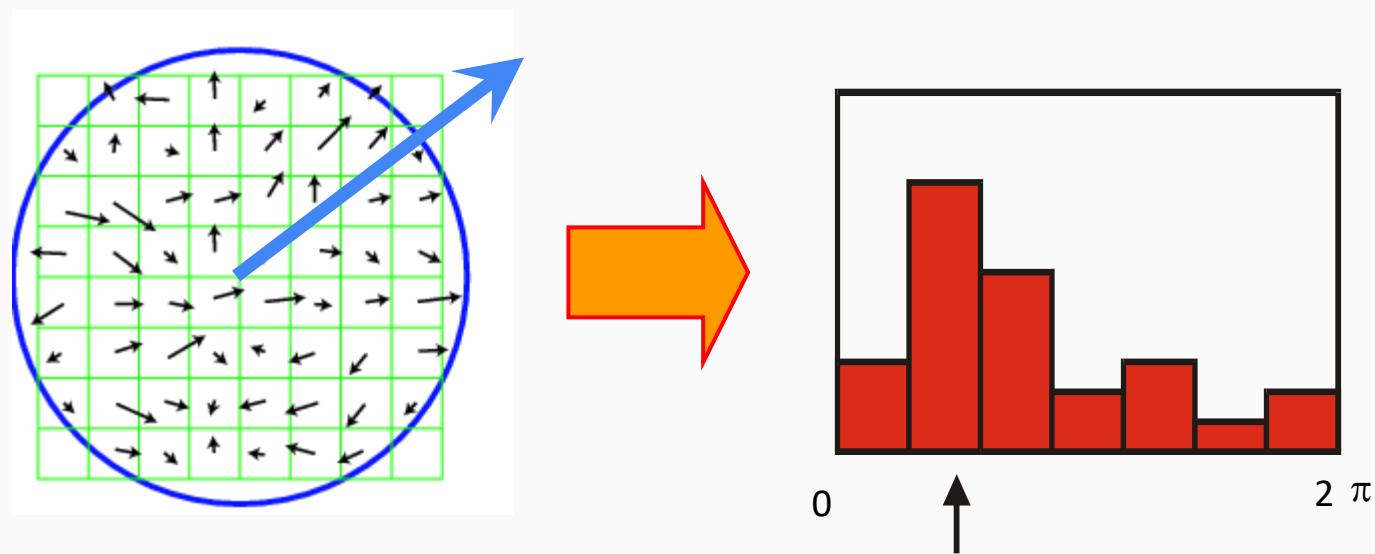
- To assign a unique orientation to circular image windows:
  - Create histogram of local gradient directions in the patch
  - Assign canonical orientation at peak of smoothed histogram



Source: S. Lazebnik

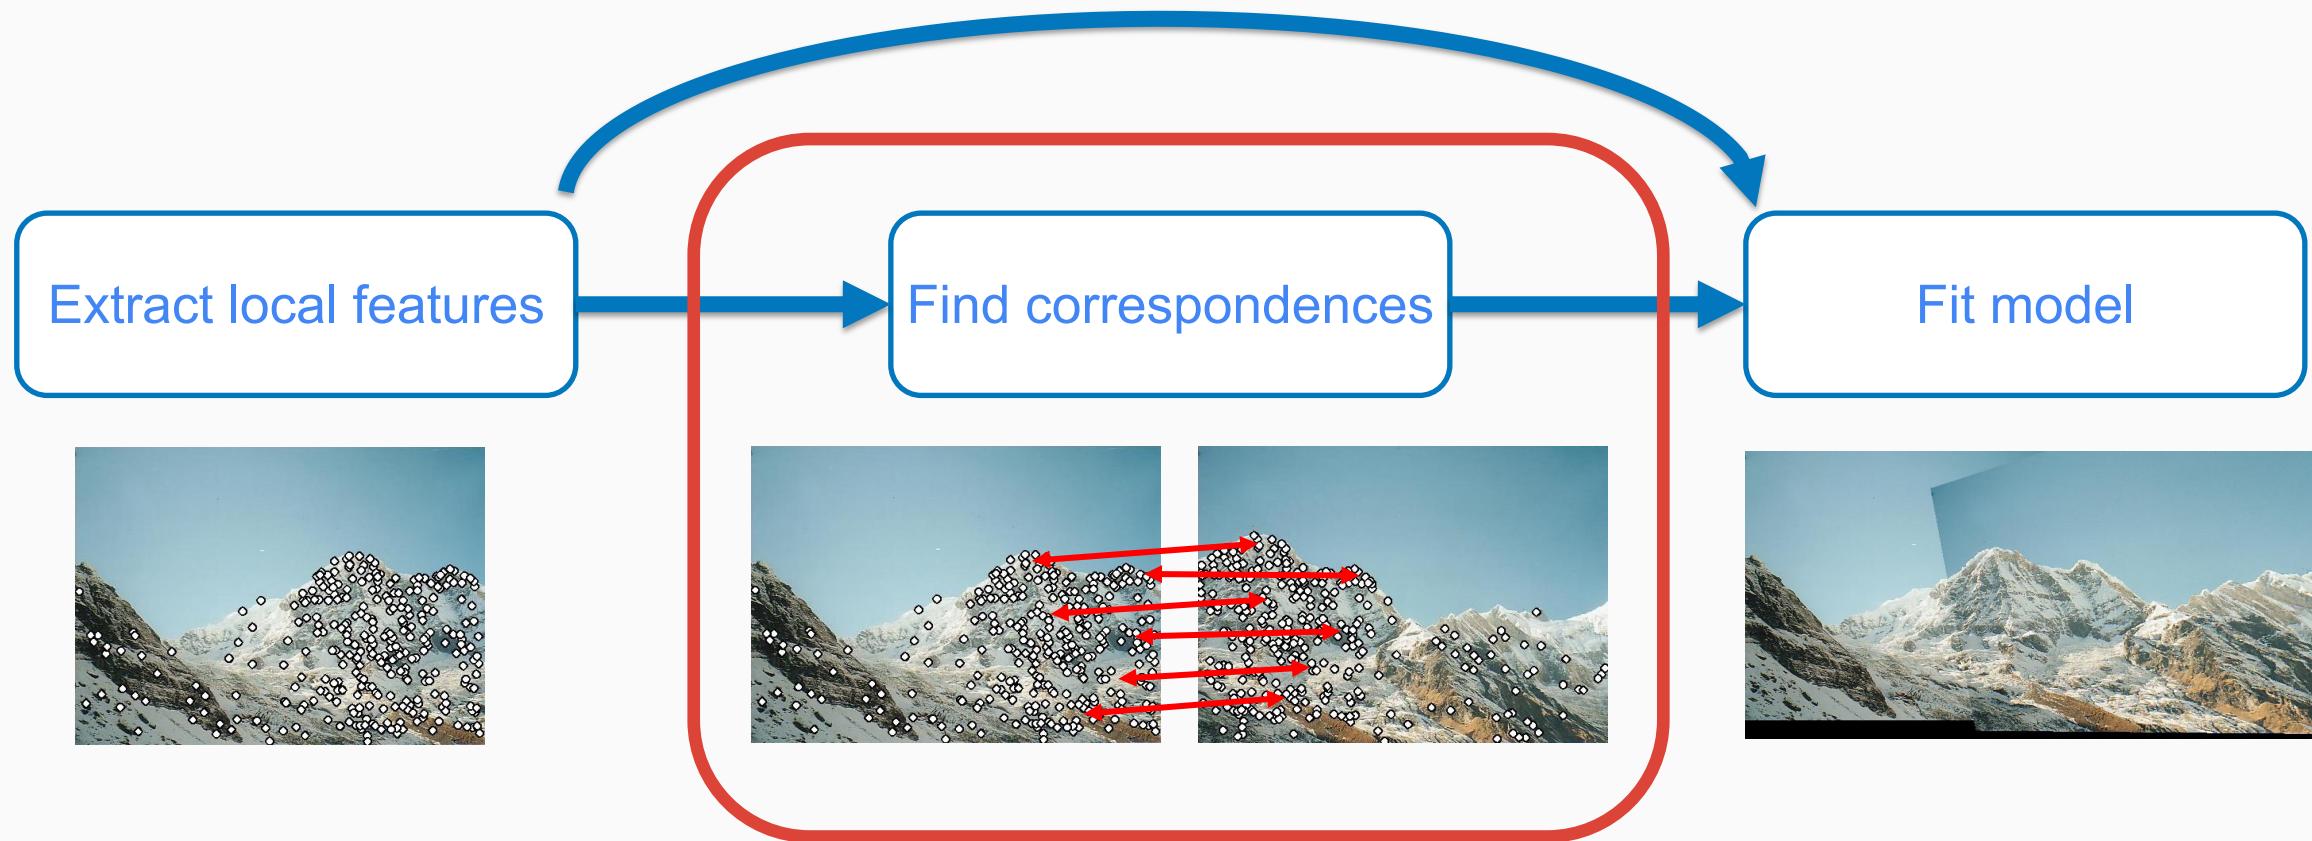
# Eliminating rotation ambiguity

- To assign a unique orientation to circular image windows:
  - Create histogram of local gradient directions in the patch
  - Assign canonical orientation at peak of smoothed histogram



Source: S. Lazebnik

# Usual steps in local feature extraction and matching



# Local feature descriptors – SIFT cont.

# SIFT features

- Detected features with characteristic scales and orientations:

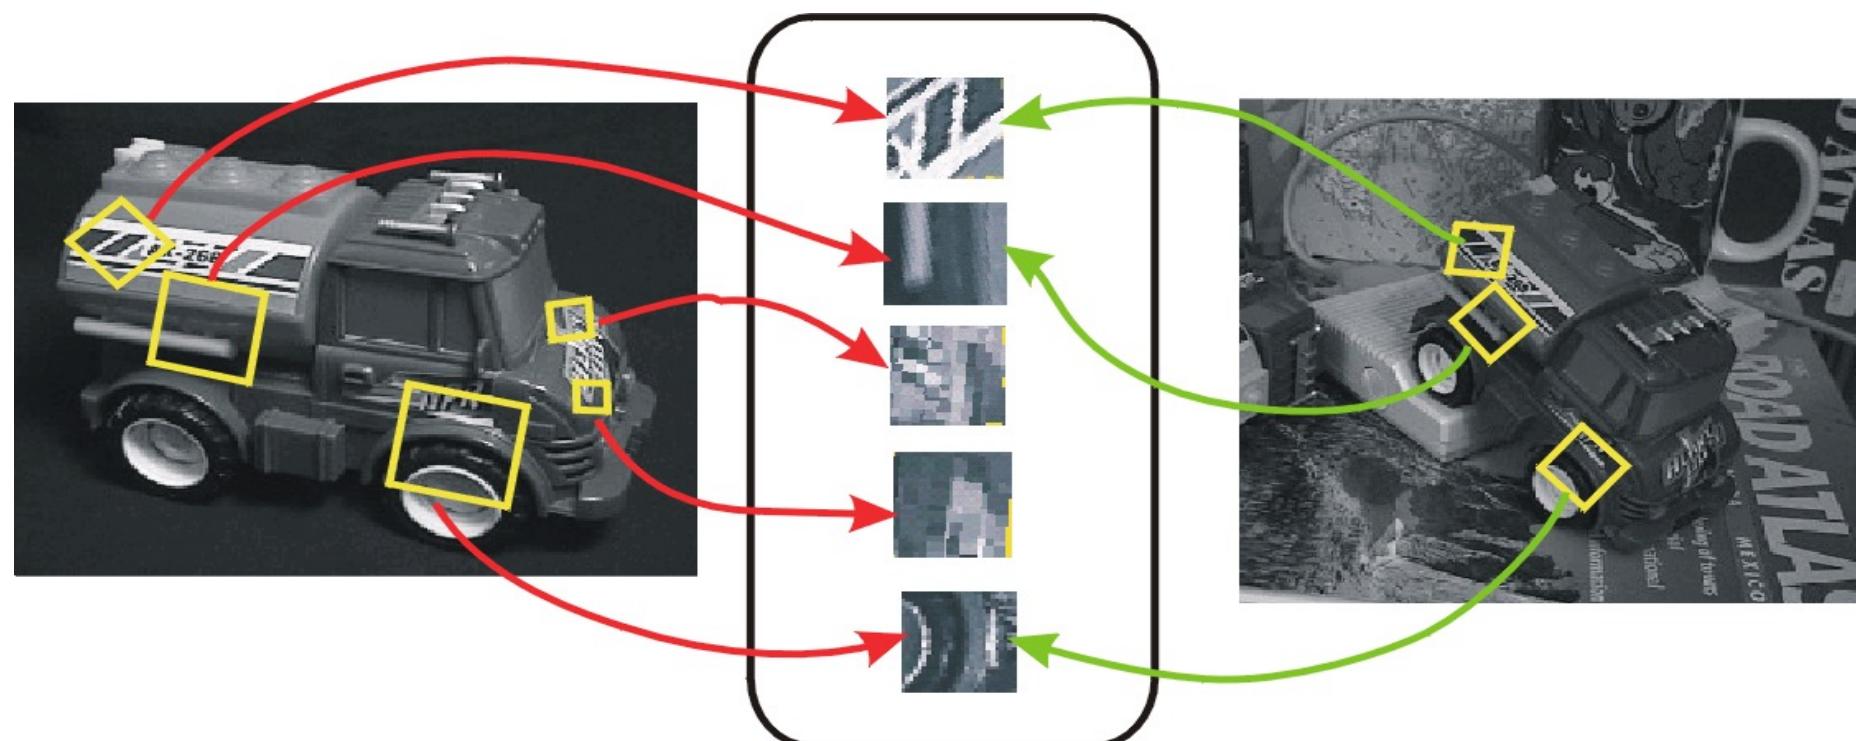


David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

Source: S. Lazebnik

# From feature detection to feature description

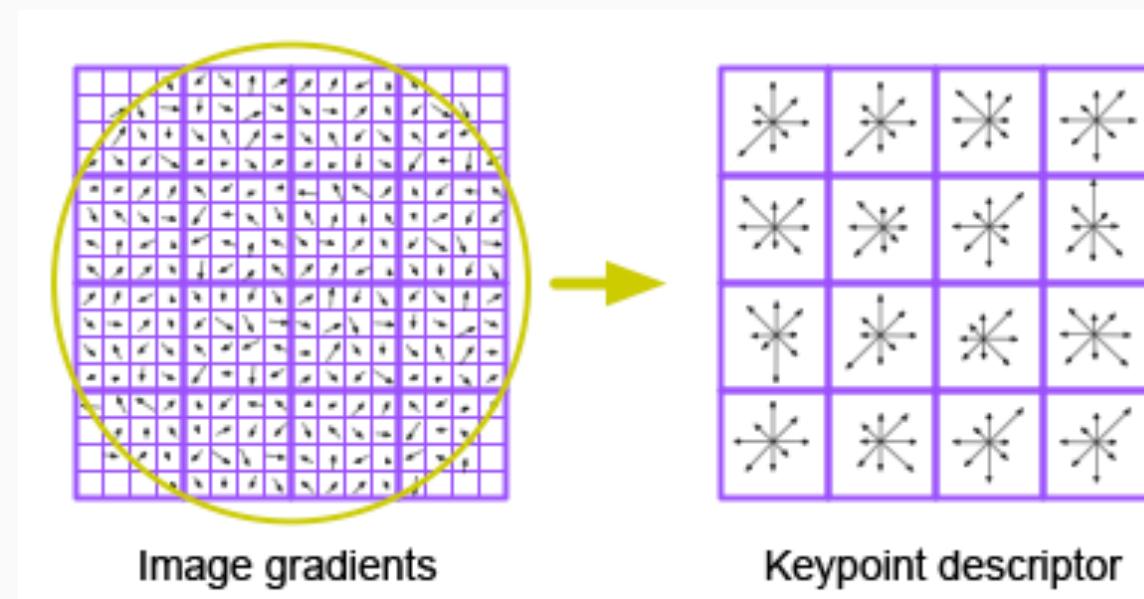
- Detection is covariant:
  - $\text{features}(\text{transform}(\text{image})) = \text{transform}(\text{features}(\text{image}))$
- Description is invariant:
  - $\text{features}(\text{transform}(\text{image})) = \text{features}(\text{image})$



Source: S. Lazebnik

# SIFT descriptors

- Inspiration: complex neurons in the primary visual cortex



D. Lowe. [\*\*Distinctive image features from scale-invariant keypoints.\*\*](#)  
IJCV 60 (2), pp. 91-110, 2004.

Source: S. Lazebnik

# SIFT descriptor computation

- use the normalized region about the keypoint
- compute gradient magnitude and orientation at each point in the region
- weight them by a Gaussian window overlaid on the circle
- create an orientation histogram over the 4 X 4 subregions of the window
- 4 X 4 descriptors over 16 X 16 sample array were used in practice. 4 X 4 times 8 directions gives a vector of 128 values.



# Summary of SIFT feature detection and description

## 1. Scale-space extrema detection

Search over multiple scales and image locations.

## 2. Keypoint localization

Fit a model to determine location and scale.

Select keypoints based on a measure of stability.

## 3. Orientation assignment

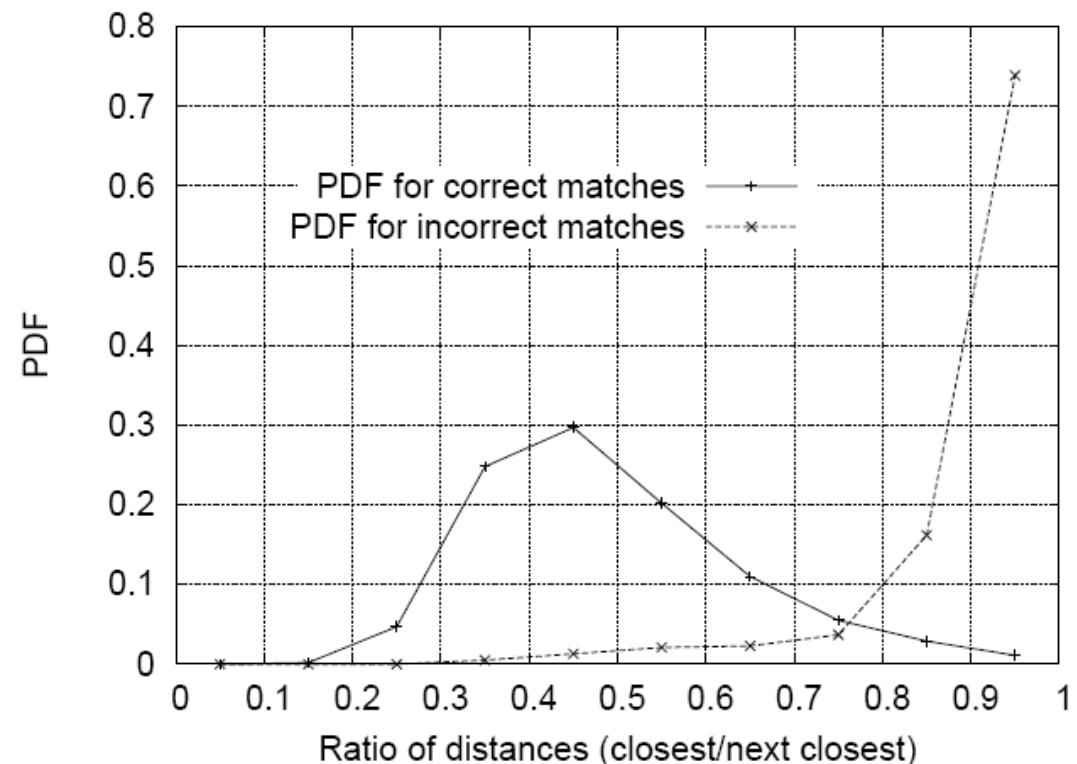
Compute best orientation(s) for each keypoint region.

## 4. Keypoint description

Use local image gradients at selected scale and rotation to describe each keypoint region.

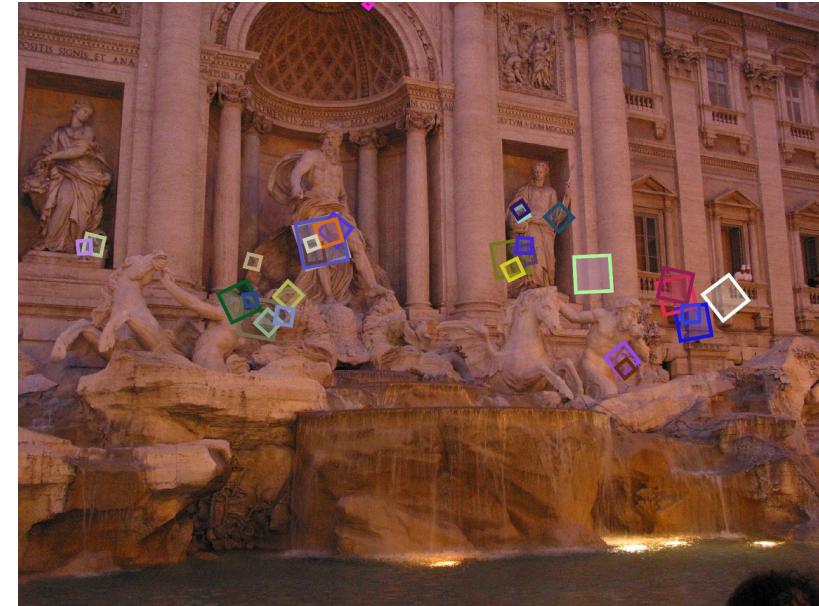
# Matching SIFT Descriptors

- Nearest neighbor (Euclidean distance)
- Threshold ratio of nearest to 2nd nearest descriptor



# Properties of SIFT

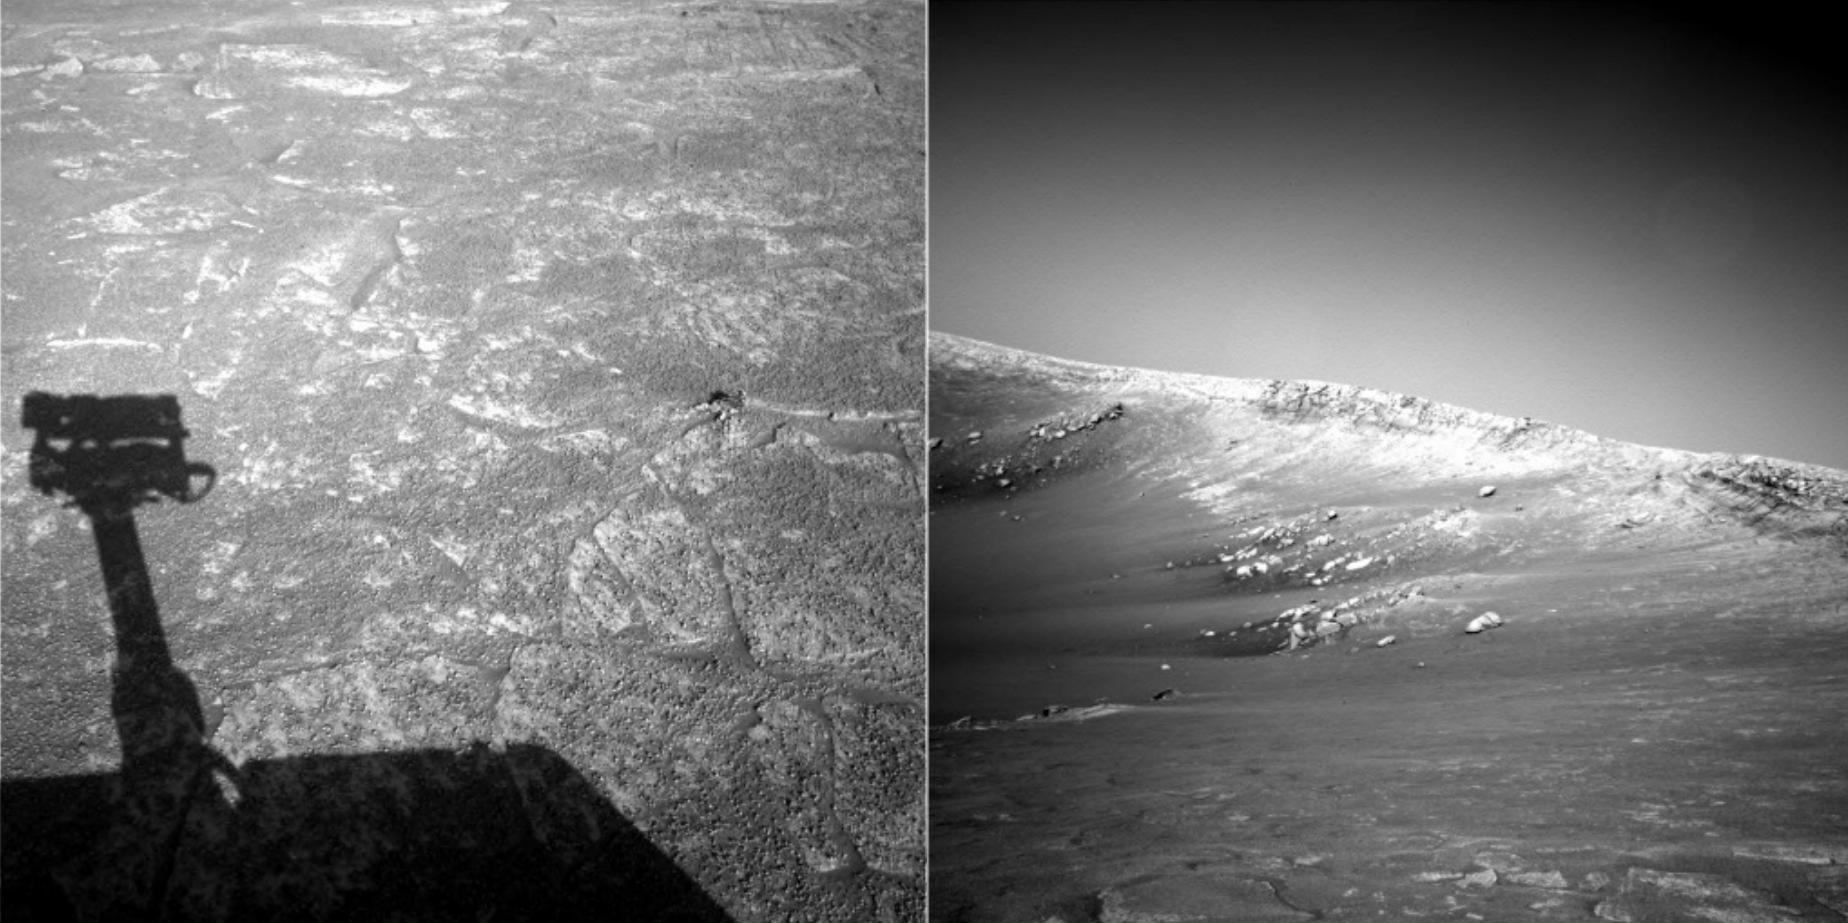
- Extraordinarily robust detection and description technique
  - Can handle changes in viewpoint  
Up to about 60 degree out-of-plane rotation
  - Can handle significant changes in illumination  
Sometimes even day vs. night
  - Fast and efficient—can run in real time
  - Lots of code available



Source: N. Snavely

# A hard keypoint matching problem

---

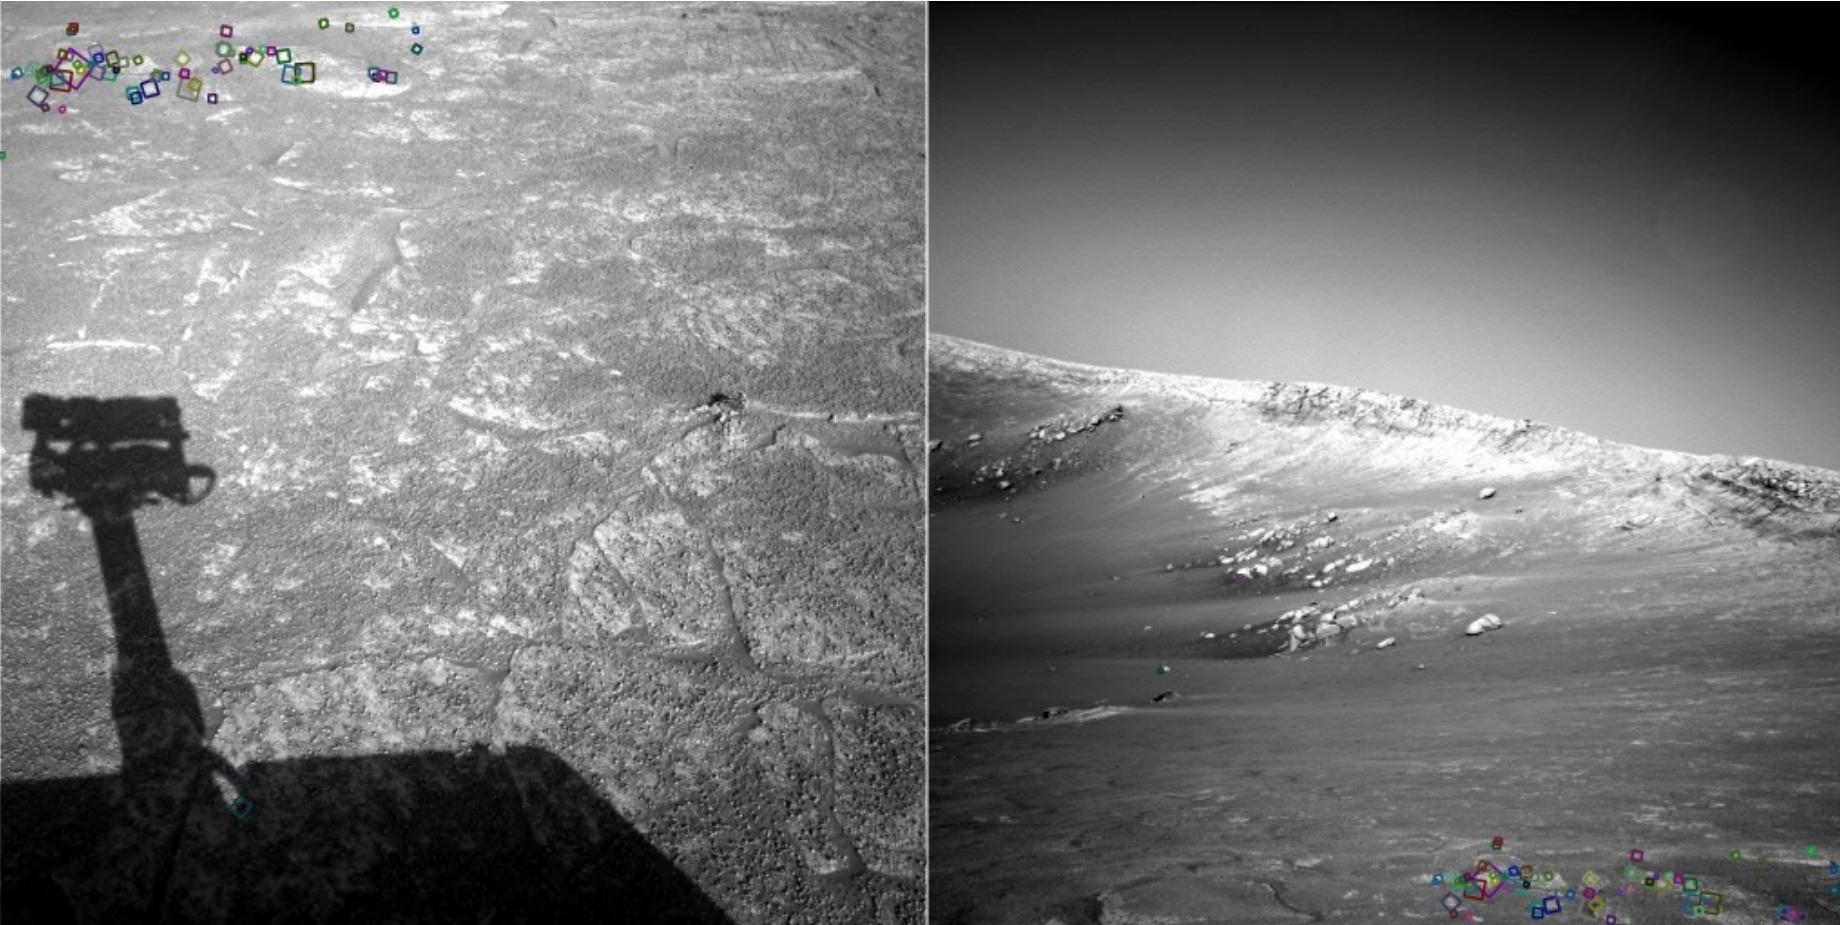


NASA Mars Rover images

Source: S. Lazebnik

Answer below (look for tiny colored squares...)

---



NASA Mars Rover images  
with SIFT feature matches  
Figure by Noah Snavely

Source: S. Lazebnik

That's all folks!